

Navigating Linux System Commands

A guide for beginners to the Shell and GNU coreutils

Sayan Ghosh

July 4, 2024

IIT Madras
BS Data Science and Applications

Disclaimer

This document is a companion activity book for the System Commands (BSSE2001) course taught by **Prof. Gandham Phanikumar** at **IIT Madras BS Program**. This book contains resources, references, questions and solutions to some common questions on Linux commands, shell scripting, grep, sed, awk, and other system commands.

This was prepared with the help and guidance of the course instructors:

Santhana Krishnan and **Sushil Pachpinde**

Copyright

© This book is released under the public domain, meaning it is freely available for use and distribution without restriction. However, while the content itself is not subject to copyright, it is requested that proper attribution be given if any part of this book is quoted or referenced. This ensures recognition of the original authorship and helps maintain transparency in the dissemination of information.

Colophon

This document was typeset with the help of **KOMA-Script** and **L^AT_EX** using the **kaobook** class.

The source code of this book is available at:

<https://github.com/sayan01/se2001-book>

(You are welcome to contribute!)

Edition

Compiled on July 4, 2024

UNIX is basically a simple operating system, but you have to be a genius to understand the simplicity.

– Dennis Ritchie

Preface

Through this work I have tried to make learning and understanding the basics of Linux fun and easy. I have tried to make the book as practical as possible, with many examples and exercises. The structure of the book follows the structure of the course *BSSE2001 - System Commands*, taught by **Prof. Gandham Phanikumar** at **IIT Madras BS Program**. .

The book takes inspiration from the previous works done for the course,

- ▶ Sanjay Kumar's Github Repository
- ▶ Cherian George's Github Repository
- ▶ Prabuddh Mathur's TA Sessions

as well as external resources like:

- ▶ Robert Elder's Blogs and Videos
- ▶ Aalto University, Finland's Scientific Computing - Linux Shell Crash Course

The book covers basic commands, their motivation, use cases, and examples. The book also covers some advanced topics like shell scripting, regular expressions, and text processing using sed and awk.

This is not a substitute for the course, but a companion to it. The book is a work in progress and any contribution is welcome at <https://github.com/sayan01/se2001-book>

Sayan Ghosh

Contents

Preface	v
Contents	vii
1 Streams, Redirections, Piping	1
1.1 Multiple Commands in a Single Line	1
1.1.1 Conjunction and Disjunction	1
1.2 Streams	4
1.3 Redirection	5
1.3.1 Standard Output Redirection	5
1.3.2 Standard Error Redirection	7
1.3.3 Appending to a File	9
1.3.4 Standard Input Redirection	10
1.3.5 Here Documents	11
1.3.6 Here Strings	12
1.4 Pipes	12
1.4.1 UNIX Philosophy	13
1.4.2 Multiple Pipes	13
1.4.3 Piping Standard Error	18
1.4.4 Piping to and From Special Files	18
1.4.5 Named Pipes	19
1.4.6 Tee Command	21
1.5 Command Substitution	22
1.6 Arithmetic Expansion	23
1.6.1 Using variables in arithmetic expansion	23
1.7 Process Substitution	23
1.8 Summary	26

List of Figures

1.1	Standard Streams	4
1.2	File Descriptor Table	4
1.3	Pipes	12

List of Tables

1.1	Pipes, Streams, and Redirection syntax	26
-----	--	----

1.1 Multiple Commands in a Single Line

Sometimes we may want to run multiple commands in a single line. For example, we may want to run two commands `ls` and `wc` in a single line. We can do this by separating the commands with a semicolon. This helps us see the output of both the commands without having the prompt in between.

For example, the following command will run `ls` and `wc` in a single line.

```
1 $ ls; wc .bashrc
2 docs down music pics programs scripts tmp vids
3 340 1255 11238 .bashrc
```

In this way of executing commands, the success or failure of one command does not affect the other command. Concisely, the commands are executed independently and sequentially. Even if a command fails, the next command will be executed.

```
1 $ date; ls /nonexistent ; wc .bashrc
2 Wed Jul 3 06:54:45 PM IST 2024
3 ls: cannot access '/nonexistent': No such file or directory
4 340 1255 11238 .bashrc
```

1.1.1 Conjunction and Disjunction

We can also run multiple commands in a single line using conjunction and disjunction. The conjunction operator `&&` is used to run the second command only if the first command is successful. The disjunction operator `||` is used to run the second command only if the first command fails. In computer science, these operators are also known as short-circuit logical **AND** and **OR** operators.¹

```
1 $ ls /nonexistent && echo "ls successful"
2 ls: cannot access '/nonexistent': No such file or directory
3 $ ls /home && echo "ls successful"
4 lost+found sayan test1
5 ls successful
```

In the first command, the `ls` command fails, so the `echo` command is not executed. In the second command, the `ls` command is successful, so the `echo` command is executed.

The success or failure of a command is determined by the exit status of the command. If the exit status is 0, the command is successful. If the exit status is non-zero, the command is considered to have failed.

The exit status of the last command can be accessed using the special variable `?`. This variable contains the exit status of the last command executed.

1.1 Multiple Commands in a Single Line	1
1.1.1 Conjunction and Disjunction	1
1.2 Streams	4
1.3 Redirection	5
1.3.1 Standard Output Redirection	5
1.3.2 Standard Error Redirection	7
1.3.3 Appending to a File	9
1.3.4 Standard Input Redirection	10
1.3.5 Here Documents	11
1.3.6 Here Strings	12
1.4 Pipes	12
1.4.1 UNIX Philosophy	13
1.4.2 Multiple Pipes	13
1.4.3 Piping Standard Error	18
1.4.4 Piping to and From Special Files	18
1.4.5 Named Pipes	19
1.4.6 Tee Command	21
1.5 Command Substitution	22
1.6 Arithmetic Expansion	23
1.6.1 Using variables in arithmetic expansion	23
1.7 Process Substitution	23
1.8 Summary	26

1: A short-circuit logical **AND** operator returns **true** if both the operands are **true**. If the first operand is **false**, it does not evaluate the second operand and returns **false** directly.

```

1 $ ls /nonexistent
2 ls: cannot access '/nonexistent': No such file or directory
3 $ echo $?
4 2

```

Here, the exit status of the `ls` command is 2, because the file `/nonexistent` does not exist.

```

1 $ ls /home
2 lost+found  sayan  test1
3 $ echo $?
4 0

```

Here, the exit status of the `ls` command is 0, because the directory `/home` exists, and the command is successful.

Similarly, we can use the disjunction operator `||` to run the second command only if the first command fails.

```

1 $ ls /nonexistent || echo "ls failed"
2 ls: cannot access '/nonexistent': No such file or directory
3 ls failed

```

In this case, the `ls` command fails, so the `echo` command is executed. However, since the disjunction operator is a short-circuit operator, the `echo` command is executed only if the `ls` command fails. If the `ls` command is successful, the `echo` command is not executed.

```

1 $ ls /home || echo "ls failed"
2 lost+found  sayan  test1

```

In this case, the `ls` command is successful, so the `echo` command is not executed.

We can also chain multiple commands using conjunction and disjunction.

```

1
2 $ date && ls /hello || echo "echo"
3 Thu Jul  4 06:53:08 AM IST 2024
4 ls: cannot access '/hello': No such file or directory
5 echo

```

In this case, the `date` command is successful, so the `ls` is executed. The `ls` command fails, so the `echo` command is executed.

However, even if the first command fails, and the second command is skipped, the third command is executed.

```

1 $ ls /hello && date || echo echo
2 ls: cannot access '/hello': No such file or directory
3 echo

```

In this case, the `ls` command fails, so the `date` command is not executed. However, the `echo` command is executed as the exit status of the `ls` command is non-zero.

To make the `echo` command execute only if the `date` command is run and fails, we can use parentheses.

```

1 $ ls /hello && (date || echo echo)
2 ls: cannot access '/hello': No such file or directory

```

Although the parentheses look like grouping a mathematical expression, they are actually a subshell. The commands inside the parentheses are executed in a subshell. The exit status of the subshell is the exit status of the last command executed in the subshell.

We can see the subshell count using the `echo $BASH_SUBSHELL`.

```
1 $ echo $BASH_SUBSHELL
2 0
3 $ (echo $BASH_SUBSHELL)
4 1
5 $ ( :;(echo $BASH_SUBSHELL ))
6 2
7 $ ( :(:;(echo $BASH_SUBSHELL )) )
8 3
```

Remark 1.1.1 When nesting in more than one subshell, simply putting two parentheses side by side will not work. This is because the shell will interpret that as the mathematical evaluation of the expression inside the parentheses. To avoid this, we can use a colon `:` no-op command followed by a semicolon `;` to separate the two parentheses.

Remark 1.1.2 Setting up an environment takes up time and resources. Thus, it is better to avoid creating subshells unless necessary.

1.2 Streams

There are three standard streams in Unix-like operating systems:

1. **Standard Input (stdin):** This is the stream where the input is read from. By default, the standard input is the keyboard.
2. **Standard Output (stdout):** This is the stream where the output is written to. By default, the standard output is the terminal.
3. **Standard Error (stderr):** This is the stream where the error messages are written to. By default, the standard error is the terminal.

There are also other numbered streams, such as 3, 4, etc., which can be used to read from or write to files.

However, sometimes a process may need to take input from a file or send output to a file. When this is required, the standard stream is mapped to a file. This is known as **redirection**.

To maintain which file or resource is mapped to which stream, the operating system maintains a table known as the **file descriptor table**. The file descriptor table is a table that maps the file descriptors to the files or resources.

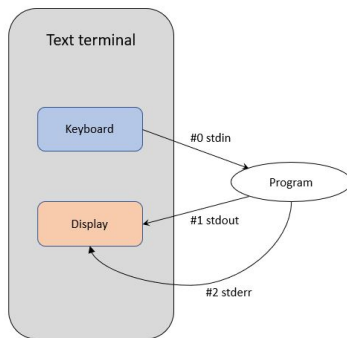


Figure 1.1: Standard Streams

Definition 1.2.1 (File Descriptor) A **file descriptor** is a process-unique identifier that the operating system assigns to a file or resource. The file descriptor is an integer that is used to identify the file or resource.

The default file descriptors are:

1. **Standard Input (stdin):** File descriptor 0
2. **Standard Output (stdout):** File descriptor 1
3. **Standard Error (stderr):** File descriptor 2

2: We have covered inodes and inode tables in detail in the Chapter ?? chapter.




Figure 1.2: File Descriptor Table

In the traditional implementation of Unix, file descriptors index into a per-process file descriptor table maintained by the kernel, that in turn indexes into a system-wide table of files opened by all processes, called the file table. This table records the mode with which the file (or other resource) has been opened: for reading, writing, appending, and possibly other modes. It also indexes into a third table called the inode table² that describes the actual underlying files. To perform input or output, the process passes the file descriptor to the kernel through a system call, and the kernel will access the file on behalf of the process. The process does not have direct access to the file or inode tables.

The file descriptors of a process can be inspected in the `/proc` directory. The `/proc` directory is a pseudo-filesystem that provides an interface to kernel data structures. The `/proc/PID/fd` directory contains the file descriptors of the process with the PID.

Most GNU core utilities that accept a file as an argument will also work without the argument and will read from the standard input. This behaviour lets us chain commands together using pipes easily without any explicit file handling.

```

1 $ cat
2 hello
3 hello
4 This command is repeating the input
5 This command is repeating the input
6 Press Ctrl+D to exit
7 Press Ctrl+D to exit

```

The `cat` command is usually used to print the contents of one or more files. However, when no file is provided, it reads from the standard input. In this case, the standard input is the keyboard. This is very useful when we want to repeat the input or when we want to read from the standard input.

Question 1.2.1 Can we also use `cat` to write to a file?

Answer 1.2.1 Yes, we can use the `cat` command to write to a file. When no file is provided, the `cat` command reads from the standard input. So the input will be printed back to the standard output. However, if we can somehow change the standard output to a file, then the input will be written to the file.

1.3 Redirection

Redirection is the process of changing the standard streams of a process. This is done by the shell before the process is executed. The shell uses the `<`, `>`, and `2>` operators to redirect the standard input, standard output, and standard error streams of a process.

As the shell is responsible for redirection, the process is not aware of the redirection. The process reads from or writes to the file descriptor it is given by the shell. The process does not know whether the file descriptor is a file, a terminal, or a pipe.

However, there are ways for a process to guess whether the file descriptor is a terminal or a file. This is done by using the `isatty()` function. The `isatty()` function returns 1 if the file descriptor is a terminal, and 0 if the file descriptor is a file.

1.3.1 Standard Output Redirection

The standard output of a process can be redirected to a file using the `>` operator. Observe the following example.

```

1 $ date > date.txt
2 $ cat date.txt
3 Thu Jul  4 07:37:27 AM IST 2024

```

Here, the output of the `date` command is redirected to the file `date.txt`. The `cat` command is then used to print the contents of the file `date.txt`.

The process did not print the output to the terminal. Instead, the shell changed the standard output of the process to the file `date.txt`.

We can also use redirection to create a file. If the file does not exist, the shell will create the file. If the file exists, the shell will truncate the file.

```
1 | $ > empty.txt
```

This command creates an empty file `empty.txt`. The `>` operator is used to redirect the output of the command to the file `empty.txt`. Since there is no output, the file is empty.

We can also use redirection along with the `echo` command to create a file with some content.

```
1 | $ echo "hello" > hello.txt
2 | $ cat hello.txt
3 | hello
```

Here, the output of the `echo` command is redirected to the file `hello.txt`. The `cat` command is then used to print the contents of the file `hello.txt`.

Recall that the `cat` command will read from the standard input if no file is provided. We can use this to create a file with the input from the standard input.

```
1 | $ cat > file.txt
2 | hello, this is input typed from the keyboard
3 | I am typing more input
4 | all this is being written to the file
5 | to stop, press Ctrl+D
6 | $ cat file.txt
7 | hello, this is input typed from the keyboard
8 | I am typing more input
9 | all this is being written to the file
10 | to stop, press Ctrl+D
```

It is important to note that the process of creating the file if it does not exist, and truncating the file if it exists, along with redirecting the output, is done by the shell, not the process. All of this is done before the process is executed. This creates an interesting exemplar when using redirection with the `ls` command.

Output contains output

```
1 | $ ls
2 | hello
3 | $ ls > output.txt
4 | $ cat output.txt
5 | hello
6 | output.txt
```

Since the file is created even before the process is executed, the file is also a part of the output of the `ls` command. This is why the file `output.txt` is also printed by the `cat` command.

Output format changes

Another interesting example using the `ls` command is when the directory contains more than one file. The output of `ls` is usually formatted to fit the terminal. However, when the output is redirected to a file, the output is not in a column format. Instead, each file is printed on a new line.


```

1 $ ls
2 hello  output.txt
3 $ ls > output.txt
4 $ cat output.txt
5 hello
6 output.txt

```

Observe that the output of the `ls` command is not in a column when redirected to a file. But how does `ls` know that the output is not a terminal? It uses the file descriptor to guess whether the output is a terminal or a file. If the file descriptor is a terminal, then the output is formatted to fit the terminal. If the file descriptor is a file, then the output is not formatted.

However, we can always force the output to be single-column using the `-l` option.

```

1 $ ls -l
2 hello
3 output.txt

```

This behaviour is not limited to the `ls` command. Most commands usually strip out the formatting when the output is redirected to a file.

If you are using a terminal that supports ANSI escape codes³, you can use the `ls` command with the `-color=auto` option to get colored output. However, when the output is redirected to a file, the ANSI escape codes are not printed. This is because the ANSI escape codes are not printable characters. They are control characters that tell the terminal to change the color of the text.

3: ANSI escape codes are special sequences of characters that are used to control the terminal. For example, to change the color of the text, the ANSI escape code `\033[31m` is used. Similarly other ANSI escape codes are used to change the text style, and the position and state of the cursor.

Output and Input from same file

If you try to redirect the output of a command to a file, and then also try to read from the same file, you will get an empty file. This is because the file is truncated before the command is executed.

```

1 $ echo "hello" > output.txt
2 $ cat output.txt
3 hello
4 $ cat output.txt > output.txt
5 $ cat output.txt
6 $

```

Remark 1.3.1 Although we can simply use `>` to redirect the output to a file, the full syntax is `1>`. This is because the file descriptor for the standard output is 1. The `1>` is used to redirect the standard output to a file. However, since the standard output is the default output, we can omit the 1 and use only `>`.

1.3.2 Standard Error Redirection

The standard error of a process can be redirected to a file using the `2>` operator. Observe the following example.

```

1 | $ ls /nonexistent 2> error.txt
2 | $ cat error.txt
3 | ls: cannot access '/nonexistent': No such file or directory

```

Here, the error message of the `ls` command is redirected to the file `error.txt`. The `cat` command is then used to print the contents of the file `error.txt`.

It is important to realise that each process has two streams to output to, the standard output and the standard error. The standard output is usually used to print the output of the process, while the standard error is used to print the error messages.

This helps us differentiate between the output and the error messages. Also, if the output of a process is redirected to a file, the error will still be printed to the terminal. This is because the standard error is not redirected to the file. This makes debugging easier, as the error messages are not lost.

```

1 | $ ls -d /home /nonexistent > output.txt
2 | ls: cannot access '/nonexistent': No such file or directory
3 | $ cat output.txt
4 | /home

```

Here, the output of the `ls` command is redirected to the file `output.txt`. However, the error message is still printed to the terminal.

We can redirect both the standard output and the standard error to files using the `>` and `2>` operators.

```

1 | $ ls -d /home /nonexistent > output.txt 2> error.txt
2 | $ cat output.txt
3 | /home
4 | $ cat error.txt
5 | ls: cannot access '/nonexistent': No such file or directory

```

Redirecting both streams to the same file

Lets try to redirect both the standard output and the standard error to the same file.

```

1 | $ ls -d /home /nonexistent > output.txt 2> output.txt
2 | $ cat output.txt
3 | /home
4 | nnot access '/nonexistent': No such file or directory

```

Why did the error message get mangled? This is because the shell truncates the file before the process is executed. So the error is written to the file, and then the output message is written to the same file, overwriting the error partially. Observe that only the first six characters of the error message are mangled, the same size of the output.⁴

The correct way to redirect both the standard output and the standard error to the same file is to use the `2>&1` operator. This means that the standard error is redirected to the standard output. Here, the `1` is the file descriptor for the standard output. The `&` is used to tell the shell that the `1` is a file descriptor, not a file.

```

1 | $ ls -d /home /nonexistent > output.txt 2>&1
2 | $ cat output.txt
3 | ls: cannot access '/nonexistent': No such file or directory

```

4: The output message is `/home`, followed by a newline character. This makes the output message 6 characters long. The shell first writes the error message to the file (because that is printed first by the `ls` command), and then overwrites the first 6 bytes of the file with the output message. Since the 6th byte is a newline character, it looks like there are two lines in the file.

```
4 | /home
```

However, the order is important. The `2>&1` operator should be placed at the end of the command. If it is placed at the beginning, then the standard error will be redirected to the standard output, which, at that point, is the terminal. Then the standard output will be redirected to the file. Thus, only the standard output will be redirected to the file.

```
1 | $ ls -d /home /nonexistent 2>&1 > output.txt
2 | ls: cannot access '/nonexistent': No such file or directory
3 | $ cat output.txt
4 | /home
```

1.3.3 Appending to a File

The `>` operator is used to redirect the output to a file. If the file does not exist, the shell will create the file. If the file exists, the shell will truncate the file. However, if we want to append the output to the file, we can use the `>>` operator.

```
1 | $ echo "hello" > hello.txt
2 | $ cat hello.txt
3 | hello
4 | $ echo "world" > hello.txt
5 | $ cat hello.txt
6 | world
7 | $ echo "hello" > hello.txt
8 | $ echo "world" >> hello.txt
9 | $ cat hello.txt
10 | hello
11 | world
```

Observe that the `>` operator truncates the file, while the `>>` operator appends to the file.

We can also append the standard error to a file using the `2>>`

```
1 | $ ls /nonexistent 2> error.txt
2 | $ cat error.txt
3 | ls: cannot access '/nonexistent': No such file or directory
4 | $ daet 2>> error.txt
5 | $ cat error.txt
6 | ls: cannot access '/nonexistent': No such file or directory
7 | bash: daet: command not found
```

This is useful when we want to append the error messages to a file, like in a log file.

Circular Redirection

If we try to redirect the output of a command to the same file that we are reading from, we will get an empty file. However, if we append the output to the file, we will get an infinite loop of the output. This should not be done, as it will fill up the disk space. However, GNU core utilities `cat` is smart enough to detect this and will not read from the file.

```
1 | $ echo hello > hello
2 | $ cat hello >> hello
3 | cat: hello: input file is output file
```

```
4 $ cat < hello > hello
5 cat: -: input file is output file
```

5: We will cover pipes in the next section.

However, it can still be tricked by using a pipe. ⁵

```
1 $ echo hello > hello
2 $ cat hello | cat >> hello
3 $ cat hello
4 hello
5 hello
```

Here we don't get into an infinite loop because the first `cat` command reads from the file and writes to the pipe. The second `cat` command reads from the pipe and writes to the file. Since the file is not read from and written to at the same time, we don't get into an infinite loop.

However, BSD utilities like `cat` do not have this check, thus giving the same input and output file will result in an infinite loop.

1.3.4 Standard Input Redirection

The standard input of a process can be redirected from a file using the `<` operator. Although most commands accept a filename in the argument to read from, so we can directly provide the file in the argument. However, some commands do not accept a filename and only read from the standard input. In such cases, we can use the `<` operator to redirect the standard input from a file.

```
1 $ wc < ~/.bashrc
2 340 1255 11238
```

Although the `wc` command accepts a filename as an argument, we can also redirect the standard input from a file using the `<` to the command. However, now the output of the `wc` command is different than when we provide the filename as an argument.

```
1 $ wc ~/.bashrc
2 340 1255 11238 /home/sayan/.bashrc
```

When we provide the filename as an argument, the `wc` command prints the number of lines, words, and characters in the file, followed by the filename. However, when we redirect the standard input from a file, the `wc` command prints only the number of lines, words, and characters in the file. This is because it does not know that the input is coming from a file. For `wc`, the input is just a stream coming from the standard input.

```
1 $ wc
2 Hello, this is input from the keyboard
3 I can simply type more input here
4 All of this is taken as standard input
5 by wc command
6 Stop by pressing Ctrl+D
7      5      29     150
```

Another example where giving filename is not possible is the `read` command. The `read` command reads a line from the standard input and assigns it to a variable. The `read` command does not accept a filename as an argument. So we have to use the `<` operator to redirect the standard input from a file.

```

1 $ read myline < /etc/hostname
2 $ echo $myline
3 rex

```

Here, the `read` command reads a line from the file `/etc/hostname` and assigns it to the variable `myline`. The `echo` command is then used to print the value of the variable `myline`. To access the value of the variable, we use the `$` operator.⁶

6: This will be covered in depth in the Chapter ?? chapter.

1.3.5 Here Documents

A **here document** is a way to provide input to a command without using a file. The input is provided directly in the command itself. This is done by using the `<<` operator followed by a delimiter.

```

1 $ wc <<EOF
2 This is acting like a file
3 however the input is directly provided
4 to the shell or the script
5 this is mostly used in scripts
6 we can use any delimiter instead of EOF
7 to stop, type the delimiter on a new line
8 EOF
9      6      41      206

```

We can optionally provide a hyphen `-` after the `<<` to ignore leading tabs. This is useful when the here document is indented.

```

1 $ cat <<EOF
2 > hello
3 >     this is space
4 >     this is tab
5 > EOF
6 hello
7     this is space
8     this is tab
9 $ cat <<-EOF
10 hello
11     this is space
12     this is tab
13 EOF
14 hello
15     this is space
16 this is tab

```

The delimiter can be any string. However, it is common to use `EOF` as the delimiter. The delimiter should be on a new line.

```

1 $ wc <<END
2 If I want to have
3 EOF
4 as part of the input
5 then I can simply change the delimiter
6 END
7      4      18      82

```

Here-documents also support variable expansion. The shell will expand the variables before passing the input to the command. Thus it is really

useful in scripts if we want to provide a template which is filled with the values of the variables.

```
1 $ name="Sayan"
2 $ cat <<EOF
3 Hello, $name
4 This is a here document
5 EOF
6 Hello, Sayan
7 This is a here document
```

1.3.6 Here Strings

A **here string** is a way to provide input to a command without using a file. The input is provided directly in the command itself. This is done by using the `<<` operator followed by the input. It is very similar to the here document, but the input is provided in a single line. Due to this, a delimiter is not required.

```
1 $ wc <<< "This is a here string"
2      1      5     22
```

It can also perform variable expansion.

```
1 $ cat <<< "Hello, $USER"
2 Hello, sayan
```

Remark 1.3.2 Both heredocs and herestrings are simply syntactic sugar. They are similar to using echo and piping the output to a command. However, it requires one less process to be executed.

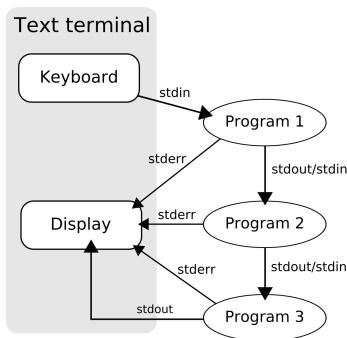


Figure 1.3: Pipes

1.4 Pipes

Pipes are the holy grail of Unix-like operating systems. They are the most important concept to understand in Unix-like operating systems.

Definition 1.4.1 (Pipe) A **pipe** is a way to connect the standard output of one process to the standard input of another process. This is done by using the `|` operator.

Think of the shell as a factory, and the commands as machines in the factory. Pipes are like conveyor belts that connect the machines. It would be a pain if we had to manually collect the produce of one machine and then feed it to the next machine. Conveyors make it easy by automatically taking the output of one machine and feeding it to the next machine.

Pipes are the same, but for processes in shell.

```
1 $ date
2 Thu Jul  4 09:05:30 AM IST 2024
3 $ date | wc
4      1      7     32
```

1.4.1 UNIX Philosophy

Each process simply takes the standard input and writes to the standard output. How those streams are connected is not the concern of the process. This way of simply doing one thing and doing it well is the Unix philosophy⁷ which says that each program should do one thing and do it well. Each process should take input from the standard input and write to the standard output. The output of the process should be easy to parse by another process. The output of the process should not contain unnecessary information. This makes it easy to chain commands together using pipes.

7: The Unix philosophy, originated by Ken Thompson, is a set of cultural norms and philosophical approaches to minimalist, modular software development. It is based on the experience of leading developers of the Unix operating system. Read more [online](#).

1.4.2 Multiple Pipes

There is no limit to the number of pipes that can be chained together. It simply means that the output of one process is fed to the input of the next process. This simple but powerful construct lets the user do any and all kinds of data processing.

Imagine you have a file which contains a lot of words, you want to find which word is present the most number of times. You can either write a program in C or Python, etc., or you can use the power of pipes to do it in a single line using simple GNU coreutils.

I have a file `alice_in_wonderland.txt` which contains the entire text of the book Alice in Wonderland. I want to find the words that are present the most number of times.

There are some basic preprocessing you would do even if you were writing a program. You would convert all the words to lowercase, and remove any punctuation. This can be done using the `tr` command. Then you would split the text into each word. This can be done using the `tr` command. Then you would find the count of each word. There are two ways of doing this, either you can use a dictionary⁸ to store the frequency of each word, and increase it as you iterate over the entire text once. Or you can first sort the words, then simply count the number of times each word is repeated. Since repeated words would always be consecutive, you can simply count the number of repeated words without having to store the frequency of each word. This can be done using the `sort` and `uniq` commands. Finally, you would sort the words based on the frequency and print the top 10 words. This can be done using the `sort` and `head` commands.

8: A dictionary in python is a hash map. It is a data structure that maps keys to values. The keys are unique, and the values can be accessed using the keys. It has amortized constant time complexity for insertion, deletion, and lookup.

Lets see how we actually do this using pipes.

```

1 $ ls alice_in_wonderland.txt
2 alice_in_wonderland.txt
3 $ tr 'A-Z' 'a-z' < alice_in_wonderland.txt | tr -cd 'a-z ' | tr '
   ' '\n' | grep . | sort | uniq -c | sort -nr | head
4     1442 the
5         713 and
6         647 to
7         558 a
8         472 she
9         463 it
10        455 of

```

```

11 | 435 said
12 | 357 i
13 | 348 alice

```

Let's go over each command one by one and see what it does.

tr 'A-Z' 'a-z'

tr is a command that translates characters. The first argument is the set of characters to be translated, and the second argument is the set of characters to translate to. Here, we are translating all uppercase letters to lowercase. This command converts all uppercase letters to lowercase. This is done to make the words case-insensitive.

```

1 | $ tr 'A-Z' 'a-z' < alice_in_wonderland.txt | head -n20
2 | alice's adventures in wonderland
3 |
4 |         alice's adventures in wonderland
5 |
6 |             lewis carroll
7 |
8 |             the millennium fulcrum edition 3.0
9 |
10 |
11 |
12 |
13 |             chapter i
14 |
15 |             down the rabbit-hole
16 |
17 |
18 |     alice was beginning to get very tired of sitting by her sister
19 | on the bank, and of having nothing to do: once or twice she had
20 | peeped into the book her sister was reading, but it had no
21 | pictures or conversations in it, 'and what is the use of a book,'

```

Remark 1.4.1 Since the text is really big, I am going to filter all the intermediate outputs also through the head command.

tr -cd 'a-z'

The -c option is used to complement the set of characters. The -d option is used to delete the characters. Here, we are telling the tr command to delete all characters except lowercase letters and spaces. This is done to remove all punctuation and special characters. Observe the difference in output now.

```

1 | $ tr 'A-Z' 'a-z' < alice_in_wonderland.txt | tr -cd 'a-z' | head
   | -c500
2 | alices adventures in wonderland          alices adventures
   | in wonderland                          lewis carroll
   |         the millennium fulcrum edition
   |         chapter i                      down the rabbithole alice
   | was beginning to get very tired of sitting by her sisteron
   | the bank and of having nothing to do once or twice she
   | hadpeeped into the book her sister was reading but it had
   | nopictures or conversations in it and what is the use of a
   | bookthought alice w

```


Remark 1.4.2 After we have removed all the punctuation, the text is now only a single line. This is because we have removed all the characters, including the newline characters. Thus to restrict the output, I am using the `head -c500` instead of `head -n20`.

tr ' ' '\n'

The `tr` command is used to translate characters. Here, we are translating spaces to newline characters. This is done to split the text into words. We are doing this so that each word is on a new line. This is helpful as the `sort` and `uniq` commands work on lines.

```
1 $ tr 'A-Z' 'a-z' < alice_in_wonderland.txt | tr -cd 'a-z ' | tr '
  ' '\n' | head -20
2 alices
3 adventures
4 in
5 wonderland
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21 alices
```

grep .

Now we have each word on a new line. But observe that there are many empty lines. This is because of the multiple spaces between words and spaces around punctuation. We can remove these empty lines using the `grep .` command.⁹

9: We will learn more about regular expressions in the Chapter ?? chapter.

```
1 $ tr 'A-Z' 'a-z' < alice_in_wonderland.txt | tr -cd 'a-z ' | tr '
  ' '\n' | grep . | head -20
2 alices
3 adventures
4 in
5 wonderland
6 alices
7 adventures
8 in
9 wonderland
10 lewis
11 carroll
12 the
13 millennium
14 fulcrum
15 edition
```

```

16 | chapter
17 | i
18 | down
19 | the
20 | rabbithole
21 | alice

```

Now we are almost there. We have each word on a new line. Now we can pass this entire stream of words to the `sort` command, this will sort the words.

sort

Sort by default sorts the words in lexicographical order. This is useful as repeated words would be consecutive. This is important as the `uniq` command only works on consecutive lines only.

```

1 | $ tr 'A-Z' 'a-z' < alice_in_wonderland.txt | tr -cd 'a-z ' | tr '
   | ' '\n' | grep . | sort | head -20
2 | a
3 | a
4 | a
5 | a
6 | a
7 | a
8 | a
9 | a
10 | a
11 | a
12 | a
13 | a
14 | a
15 | a
16 | a
17 | a
18 | a
19 | a
20 | a
21 | a

```

If you run the command yourself without the `head` command, you can see that the words are sorted. The repeated words are consecutive. In the above output we can see that the word `a` is repeated many times. Now we can use the `uniq` command to count the number of times each word is repeated.

uniq -c

The `uniq` command is used to remove consecutive duplicate lines. However, it also can count the number of times each line is repeated using the `-c` option.

```

1 | $ tr 'A-Z' 'a-z' < alice_in_wonderland.txt | tr -cd 'a-z ' | tr '
   | ' '\n' | grep . | sort | uniq -c | head -20
2 | 558 a
3 | 1 abarrowful
4 | 1 abat
5 | 1 abidefigures
6 | 1 able
7 | 79 about

```

```

8      1 aboutamong
9      1 aboutand
10     1 aboutby
11     3 abouther
12     3 aboutit
13     1 aboutthis
14     1 abouttrying
15     2 above
16     1 abranh
17     1 absenceand
18     2 absurd
19     1 acceptance
20     2 accident
21     1 accidentally

```

Great! Now we have the count of each word. However, the words are still sorted by the word. We want to sort the words by the count of the word. This can be done using the `sort` command.

sort -nr

Sort by default sorts in lexicographical order. However, we want to sort by the count of the word. The `-n` option is used to sort the lines numerically. The `-r` option is used to sort in reverse order.

Exercise 1.4.1 Try to run the same command without the `-n` option. Observe how the sorting makes sense alphabetically, but not numerically.

```

1 $ tr 'A-Z' 'a-z' < alice_in_wonderland.txt | tr -cd 'a-z ' | tr '
   ' '\n' | grep . | sort | uniq -c | sort -nr | head -20
2 1442 the
3   713 and
4   647 to
5   558 a
6   472 she
7   463 it
8   455 of
9   435 said
10  357 i
11  348 alice
12  332 you
13  332 in
14  313 was
15  241 that
16  237 as
17  202 her
18  190 at
19  169 on
20  161 all
21  158 with

```

Finally, we have the top 20 words in the file `alice_in_wonderland.txt` along with the count of each word.

Although the above command is a single line ¹⁰, it is doing a lot of processing. Still, it is very readable. This is the power of pipes.

¹⁰: Such commands are called one-liners.

Remark 1.4.3 Usually, when we come across such one-liners, it may initially seem to complicated to understand. The key to understanding such one-liners is to break them down into smaller parts and understand them one component at a time from left to right. Feel free to execute each command separately and observe the output like we did above.

1.4.3 Piping Standard Error

Pipes are used to connect the standard output of one process to the standard input of another process. However, the standard error is not connected and remains mapped to the terminal. This is because the standard error is a separate stream from the standard output.

However, we can connect the standard error to the standard input of another process using the `2>&1` operator. This is useful when we want to process the error messages of a command.

```
1 $ ls -d /home /nonexistent | wc
2 "/nonexistent": No such file or directory (os error 2)
3      1      1      6
4 $ ls -d /home /nonexistent 2>&1 | wc
5      2     10     61
```

This is same as redirecting both the streams to a single file as demonstrated earlier.

However, there is a shorter way to do this using the `|&` syntactic sugar.

```
1 $ ls -d /home /nonexistent |& wc
2      2     10     61
```

This does the exact same thing as `2>&1` followed by the pipe.

1.4.4 Piping to and From Special Files

As we discussed earlier, there are some special files in the `/dev` directory.

`/dev/null`

The `/dev/null` file is a special file that discards all the data that is written to it. It is like a black hole. All the data that is not needed can be written to this file. Usually errors are written to the `/dev/null` file.

```
1 $ ls -d /nonexistent /home 2> /dev/null
2 /home
```

The error is not actually stored in any file, thus the storage space is not wasted.

`/dev/zero`

The `/dev/zero` file is a special file that provides an infinite stream of null bytes. This is useful when you want to provide an infinite stream of data to a process.

```

1 $ head -c1024 /dev/zero > zero.txt
2 $ ls -lh zero.txt
3 -rw-r--r-- 1 sayan sayan 1.0K Jul  4 15:04 zero.txt

```

Here we are taking the first 1024 bytes from the `/dev/zero` and writing it to the file `zero.txt`. The file `zero.txt` is 1.0K in size. This is because the `/dev/zero` file provides an infinite stream of null bytes, of which 1024 bytes are taken.

Warning 1.4.1 Make sure to always use `head` or any other limiter when working with `/dev/zero` or `/dev/random` as they are infinite streams. Forgetting this can lead to the disk being filled up. `head` with the default parameter will also not work, since it depends on presence of newline characters, which is not there in an infinite stream of zeros. That is why we are using a byte count limiter using the `-c` option. If you forget to add the limiter, you can press `Ctrl+C` as quickly as possible to stop the process and then remove the file using `rm`.

`/dev/random` and `/dev/urandom`

The `/dev/random` and `/dev/urandom` files are special files that are infinite suppliers of random bytes. The `/dev/random` file is a blocking random number generator. This means that it will block if there is not enough entropy. The `/dev/urandom` file is a non-blocking random number generator. This means that it will not block even if there is not enough entropy. Both can be used to generate random numbers.

```

1 $ head -c1024 /dev/random > random.txt
2 $ ls -lh random.txt
3 -rw-r--r-- 1 sayan sayan 1.0K Jul  4 15:04 random.txt

```

Observe that here too the file is of size 1.0K. This is because we are still taking only the first 1024 bytes from the infinite stream of random bytes. However, if we `gzip` the data, we can see that the zeros file is much smaller than the random file.

```

1 $ gzip random.txt zero.txt
2 $ ls -lh random.txt.gz zero.txt.gz
3 -rw-r--r-- 1 sayan sayan 1.1K Jul  4 15:11 random.txt.gz
4 -rw-r--r-- 1 sayan sayan   38 Jul  4 15:10 zero.txt.gz

```

The random file is 1.1K in size, while the zero file is only 38 bytes in size. This is because the random file has more entropy and thus cannot be compressed as much as the zero file.

1.4.5 Named Pipes

A **named pipe** is a special file that provides a way to connect the standard output of one process to the standard input of another process. This is done by creating a special file in the filesystem. We have already covered named pipes in the Chapter ?? chapter.

Although a pipe is faster than a named pipe, a named pipe can be used to connect processes that are not started at the same time or from the same shell. This is because a named pipe is a file in the filesystem and can be accessed by any process that has the permission to access the file.

Try out the following example. First create a named pipe using the `mkfifo` command.

```
1| $ mkfifo pipe1
```

Then run two processes, one that writes to the named pipe, another that reads from the named pipe. The order of running the processes is not important.

Terminal 1:

```
1| $ cat /etc/profile > pipe1
```

Terminal 2:

```
1| $ wc pipe1
2|    47    146    993 pipe1
```

Exercise 1.4.2 After you have created the named pipe, try changing the order of running the other two processes. Observe that whatever is run first will wait for the other process to start. This is because a named pipe is not storing the data piped to it in the filesystem. It is simply a buffer in the memory.

A named pipe is more useful over regular files when two processes want to communicate with each other. This is because a named pipe is

1. Faster than a regular file as it does not store the data in the file system.
2. Independent of the order of launching the processes. The reader can be launched first and it will still wait for the writer to send the data.
3. Works concurrently. The writer does not need to be done writing the entire data before the reader can start reading. The reader can start as soon as the writer starts writing. The faster process will simply block until the slower process catches up.

To demonstrate the last point, try running the following commands.

Ensure that a named pipe is created.

Terminal 1:

```
1| $ grep linux pipe1
```

This process is looking for lines containing the word `linux` in the file `pipe1`. Initially it will simply block. `Grep` is a command that does not wait for the entire file to be read. It starts printing output as soon as a line containing the pattern is read.

Terminal 2:

```
1| $ tree / > pipe1
```

This command is attempting to list out each and every file on the filesystem. This takes a lot of time. However, since we are using a named pipe, the `grep` command will start running as soon as the first line is written to the pipe.

You can now observe the first terminal will start printing some lines containing the word `linux` as soon as the second terminal starts writing to the pipe.

Now try the same with a regular file.

```
1 $ touch file1 # ensure its a normal file
2 $ tree / > file1
3 $ grep linux file1
```

Since this is a regular file, we cannot start reading from the file before the entire file is written. If we do that, the `grep` command will quit as soon as it catches up with the writer, it will not block and wait.

Observe that to start getting output on the screen takes a lot longer in this method. Not to mention the disk space wasted due to this.

Remark 1.4.4 Remember that when we use redirection (`>`) to write to a file, the shell truncates the file. But when we use a named pipe, the shell knows that the file is a named pipe and does not truncate the file.

1.4.6 Tee Command

The `tee` command is a command that reads from the standard standard input and writes to the standard output and to a file. It is very useful when you want to save the output of a command but also want to see the output on the terminal.

```
1 $ ls -d /home /nonexistent | tee output.txt
2 ls: cannot access '/nonexistent': No such file or directory
3 /home
4 $ cat output.txt
5 /home
```

Observe that only the standard output is written to the file, and not the standard error. This is because pipes only connect the standard output to the standard input of the next command, and the standard error remains mapped to the terminal.

Thus in the above output, although it may look like that both the standard output and the standard error are written to the terminal by the same command, it is not so.

The standard error of the `ls` command remains mapped to the terminal, and thus gets printed directly, whereas the standard output is redirected to the standard input of the `tee` command, which then prints that to the standard output (and also writes it to the file).

You can also mention multiple files to write to.

```
1 $ ls -d /home | tee output1.txt output2.txt
2 $ cat output1.txt
3 /home
4 $ cat output2.txt
5 /home
6 $ diff output1.txt output2.txt
7 $
```

Remark 1.4.5 The `diff` command is used to compare two files. If the files are the same, then the `diff` command will not print anything. If the files are different, then the `diff` command will print the lines that are different.

We can also append to the file using the `-a` option.

```
1 $ ls -d /home | tee output.txt
2 $ ls -d /etc | tee -a output.txt
3 $ cat output.txt
4 /home
5 /etc
```

1.5 Command Substitution

We have already seen that we can run multiple commands in a subshell in bash by enclosing them in parentheses.

```
1 $ (whoami; date)
2 sayan
3 Thu Jul  4 09:05:30 AM IST 2024
```

This is useful when you simply want to print the standard output of the commands to the terminal. However, what if you want to store the output of the commands in a variable? Or what if you want to pass the standard output of the commands as an argument to another command?

To do this, we use command substitution. Command substitution is a way to execute one or more processes in a subshell and then use the output of the subshell in the current shell.

There are two ways to do command substitution.

1. Using backticks 'command' - this is the legacy way of doing command substitution. It is not recommended to use this as it is difficult to read and can be confused with single quotes. It is also harder to nest.
2. Using the `$(command)` syntax - this is the recommended way of doing command substitution. It is easier to read and nest.

Throughout this book, we will use the `$(command)` syntax and not the backticks.

```
1 $ echo "Today is $(date)"
2 Today is Thu Jul  4 09:05:30 AM IST 2024
```

Here we are using the `$(date)` command substitution to get the current date and time and then using it as an argument to the `echo` command.

```
1 $ myname="$(whoami)"
2 $ mypc="$(hostname)"
3 $ echo "Hello, $myname from $mypc"
4 Hello, sayan from rex
```

We can store the output of the command in a variable and then use it later. This is useful when you want to use the output of a command multiple times.

Remark 1.5.1 Although you do not need to use the quotes with the command substitution in this case, it is always recommended to use quotes around the variable assignment, since if the output is multiword or multiline, it will throw an error.

1.6 Arithmetic Expansion

Arithmetic expansion allows the evaluation of an arithmetic expression and the substitution of the result. The format for arithmetic expansion is:

```
1 | $(( expression ))
```

This is the reason we cannot directly nest subshells without a command between the first and the second subshell.

```
1 | $ cat /dev/random | head -c$((1024*1024)) > random.txt
2 | $ ls -lh random.txt
3 | -rw-r--r-- 1 sayan sayan 1.0M Jul  4 15:04 random.txt
```

Here we are using **arithmetic expansion** to calculate the number of bytes in 1MiB¹¹ and then using it as an argument to the head command. This results in creation of a file `random.txt` of size 1MiB.

11: 1MiB = 1024KiB = 1024 × 1024 bytes
 - this is called a mebibyte
 1MB = 1000KB = 1000 × 1000 bytes - this
 is called a megabyte
 This is a very common confusion
 amongst common people. Kilo, Mega,
 Giga are SI prefixes, while Kibi, Mebi,
 Gibi are IEC prefixes.

1.6.1 Using variables in arithmetic expansion

We can also use variables in arithmetic expansion. We don't have to use the `$` operator to access the value of the variable inside the arithmetic expansion.

```
1 | $ a=10
2 | $ b=20
3 | $ echo $((a+b))
4 | 30
```

There are other ways to do arithmetic in bash, such as using the `expr` command, or using the `let` command. However, the `${}` syntax is the most recommended way to do simple arithmetic with variables in bash.

1.7 Process Substitution

Process substitution is a way to provide the output of a process as a file. This is done by using the `<(command)` syntax.

Some commands do not accept standard input. They only accept a filename as an argument. This is the exact opposite of the issue we had with the `read` command, which accepted only standard input and not a filename. There we used the `<` operator to redirect the standard input from a file.

The `diff` command is a command that compares two files and prints out differences. It does not accept standard input. If we want to compare

differences between the output of two processes, we can use process substitution.

Imagine you have two directories and you want to compare the files in the two directories. You can use the `diff` command to compare the two directories.

```
1 $ ls
2 $ dir1 dir2
3 $ ls dir1
4 file1 file2 file3
5 $ ls dir2
6 file2 file3 file4
```

We can see that the two directories have some common files (`file2` and `file3`) and some different files (`file1` in `dir1` and `file4` in `dir2`).

However, if we have a lot of files, it is difficult to see manually which files are different.

Let us first try to save the output of the two `ls` commands to files and then compare the files using `diff`.

```
1 $ ls
2 dir1 dir2
3 $ ls dir1 > dir1.txt
4 $ ls dir2 > dir2.txt
5 $ diff dir1.txt dir2.txt
6 1d0
7 < file1
8 3a3
9 > file4
```

Great! We can see that the file `file1` is present only in `dir1` and the file `file4` is present only in `dir2`. All other files are common.

However observe that we had to create two files `dir1.txt` and `dir2.txt` to store the output of the `ls` commands. This is not efficient. If the directories contained a million files, then we would have to store tens or hundreds of megabytes of data in the files.

It sounds like a job for the named pipes we learnt earlier. Lets see how easier or harder that is.

```
1 $ ls
2 dir1 dir1.txt dir2 dir2.txt
3 $ rm dir1.txt dir2.txt
4 $ mkfifo dir1.fifo dir2.fifo
5 $ ls dir1 > dir1.fifo &
6 $ ls dir2 > dir2.fifo &
7 $ diff dir1.fifo dir2.fifo
8 1d0
9 < file1
10 3a3
11 > file4
```

Et voila! We have the same output as before, but without the data actually being stored in the filesystem. The data is simply stored in the memory till the `diff` command reads it. However observe that we had to create two named pipes, and also run the `ls` processes in the background as

otherwise they would block. Also, we have to remember to delete the named pipes after using them. This is still too much hassle.

Let us remove the named pipes and the files.

```
1 | $ rm *fifo
```

Now let us see how easy it is using process substitution.

```
1 | $ diff <(ls dir1) <(ls dir2)
2 | 1d0
3 | < file1
4 | 3a3
5 | > file4
```

Amazing! We have the same output as before, but without having to initialize anything. The process substitution does all the magic of creating temporary named pipes and running the processes with the correct redirections concurrently. It then substitutes the filenames of the named pipes in the place of the process substitution.

Process substitution is also extremely useful when comparing between expected output and actual output in some evaluation of a student's scripts.¹²

We can also use process substitution to provide input to a process running in the subshell.

```
1 | tar cf >(bzip2 -c > file.tar.bz2) folder1
```

This calls `tar cf /dev/fd/?? folder1`, and `bzip2 -c > file.tar.bz2`.

Because of the `/dev/fd/<n>` system feature, the pipe between both commands does not need to be named. This can be emulated as

```
1 | mkfifo pipe
2 | bzip2 -c < pipe > file.tar.bz2 &
3 | tar cf pipe folder1
4 | rm pipe
```

Remark 1.7.1 `tar` is a command that is used to create archives. It simply puts all the files and directories in a single file. It does not perform any compression. The `c` option is used to create an archive. The `f` option is used to mention the name of the archive. The `bzip2` command is used to compress files. The `-c` option is used to write the compressed data to the standard output. The `>` operator is used to redirect the standard output to a file. We will cover `tar` and `zips` in more detail later.

That is pretty much all you need to know about pipes and redirections. To really understand and appreciate the power of pipes and redirections, you have to stop thinking imperically (like C or Python) and start thinking in streams, like a functional programming language. Once this paradigm shift happens, you will start to see the power of pipes and redirections and will be able to tackle any kind of task in the command line.

12: Try to find if this is used in the evaluation scripts of the VM Tasks!

This example is lifted from <https://tldp.org/LDP/abs/html/process-sub.html>. If you are interested in more examples of process substitution, refer the same.

1.8 Summary

Let us quickly summarize the important syntax and commands we learnt in this chapter.

Table 1.1: Pipes, Streams, and Redirection syntax

Syntax	Command	Description
;	Command Separator	Run multiple commands in a single line
&&	Logical AND	Run the second command only if the first command succeeds
	Logical OR	Run the second command only if the first command fails
>	Output Redirection	Redirect the stdout of the process to a file
»	Output Redirection (Append)	Append the stdout of the process to a file
<	Input Redirection	Redirect the stdin of the process from a file
2>	Error Redirection	Redirect the stderr of the process to a file
2>&1	Error Redirection	Redirect the stderr of the process to the stdout
«EOF	Here Document	Redirect the stdin of the process from a block of text
«<	Here String	Redirect the stdin of the process from a string
	Pipe	Connect the stdout of one process to the stdin of another process
&	Pipe Stderr	Connect the stderr of one process to the stdin of another process
\$(command)	Command Substitution	Run a command and use the output in the current shell
\$(expression)	Arithmetic Expansion	Evaluate an arithmetic expression
<(command)	Process Substitution	Provide the output of a process as a file
>(command)	Process Substitution	Provide the input to a process from a file