

Navigating Linux System Commands

A guide for beginners to the Linux Shell and GNU coreutils

Sayan Ghosh

May 20, 2024

IIT Madras
BS Data Science and Applications

Disclaimer

This document is a companion activity book for the System Commands (BSSE2001) course taught by **Prof. Gandham Phanikumar** at **IIT Madras BS Program**. This book contains resources, references, questions and solutions to some common questions on Linux commands, shell scripting, grep, sed, awk, and other system commands.

This was prepared with the help and guidance of the course instructors:

Santhana Krishnan and **Sushil Pachpinde**

Copyright

© This book is released under the public domain, meaning it is freely available for use and distribution without restriction. However, while the content itself is not subject to copyright, it is requested that proper attribution be given if any part of this book is quoted or referenced. This ensures recognition of the original authorship and helps maintain transparency in the dissemination of information.

Colophon

This document was typeset with the help of **KOMA-Script** and **L^AT_EX** using the **kaobook** class.

The source code of this book is available at:

<https://github.com/sayan01/se2001-book>

(You are welcome to contribute!)

Edition

Compiled on May 20, 2024

UNIX is basically a simple operating system, but you have to be a genius to understand the simplicity.

– Dennis Ritchie

Preface

Through this work I have tried to make learning and understanding the basics of Linux fun and easy. I have tried to make the book as practical as possible, with many examples and exercises. The structure of the book follows the structure of the course *BSSE2001 - System Commands*, taught by **Prof. Gandham Phanikumar** at **IIT Madras BS Program**. .

The book takes inspiration from the previous works done for the course,

- ▶ Sanjay Kumar's Github Repository
- ▶ Cherian George's Github Repository
- ▶ Prabuddh Mathur's TA Sessions

as well as external resources like:

- ▶ Robert Elder's Blogs and Videos
- ▶ Aalto University, Finland's Scientific Computing - Linux Shell Crash Course

The book covers basic commands, their motivation, use cases, and examples. The book also covers some advanced topics like shell scripting, regular expressions, and text processing using sed and awk.

This is not a substitute for the course, but a companion to it. The book is a work in progress and any contribution is welcome at <https://github.com/sayan01/se2001-book>

Sayan Ghosh

Contents

Preface	v
Contents	vii
1 Basic Commands in Linux	1
1.1 Introduction	1
1.1.1 What is Linux?	1
1.1.2 Desktop Environments	2
1.1.3 Window Manager	3
1.1.4 Why Linux?	4
1.1.5 What is Shell?	4
1.1.6 Shell vs Terminal	4
1.1.7 Why the Command Line?	5
1.1.8 Command Prompt	6
1.1.9 Simple Commands in GNU Core Utils	6
1.2 Navigating the File System	21
1.2.1 What is a File System?	21
1.2.2 Why do we have proc and sys directories?	22
1.2.3 Paths	23
1.2.4 Basic Commands for Navigation	23
1.3 File Permissions	24
1.4 Inodes and Links	25
1.5 System Management and Information	26
1.6 Reading and Writing Files using cat	27
1.7 Types of Files	28
1.8 Types of Commands	28
1.9 File Manipulation	29
2 Shell Variables	31
2.1 Creating Variables	31
2.2 Printing Variables to the Terminal	32
2.3 Bash Flags	33
2.4 Signals	33
2.5 Variable Manipulation	34
2.6 Arrays	36
3 Process Management	39
3.1 What are processes?	39
APPENDIX	41
Notation	43
Alphabetical Index	45

List of Figures

1.1	Linux Distributions Usage	2
1.2	Desktop Environment Usage	3
1.3	Operating System Onion Rings	4
1.4	GNU Core Utils Logo	6
1.5	ls -l Output	8
1.6	Linux Filesystem Hierarchy	21
1.7	Relative Path	23

List of Tables

1.1	Basic Commands in GNU Core Utils	7
1.2	Manual Page Sections	10
1.3	Escape Characters in echo	14
1.4	Date Format Specifiers	18
1.5	Basic Shortcuts in Terminal	20
1.6	Linux Filesystem Hierarchy	21
1.7	Linux Filesystem Directory Classification	22

Basic Commands in Linux

1

1.1 Introduction

1.1.1 What is Linux?

Definition 1.1.1 (Linux) Linux is a **kernel** that is used in many operating systems. It is open source and free to use. Linux is not an operating system unto itself, but the core component of it.

So what is **Ubuntu**? **Ubuntu** is one of the many *distributions* that use the Linux kernel. It is a complete operating system that is free to use and open source. It is based on the **Debian** distribution of Linux. There are many other *distributions* of Linux, such as:

- **Debian** - Used primarily on servers, it is known for its stability.
 - **Ubuntu** - A commercial distribution based on Debian which is popular among new users.
 - **Linux Mint** - A distribution based on Ubuntu which is known for its ease of use. It is one of the distributions recommended to new users.
 - **Pop OS** - A distribution based on Ubuntu which is known for its focus on developers, creators, and gamers.
 - and many more
- **Red Hat Enterprise Linux (RHEL)** - A commercial distribution used primarily in enterprises. It is owned by **Red Hat** and is targeted primarily to companies with their free OS paid support model.
 - **Fedora** - A community-driven distribution sponsored by **Red Hat**. It is known for its cutting-edge features and is used by developers. It remains on the upstream of **RHEL**, receiving new features before **RHEL**.
 - **CentOS** - A discontinued distribution based on **RHEL**. It was known for its stability and was used in servers. It was downstream from **RHEL**.
 - **CentOS Stream** - It is a midstream between the upstream development in Fedora Linux and the downstream development for Red Hat Enterprise Linux.
 - **Rocky Linux** - A distribution created by the **Rocky Enterprise Software Foundation** after the announcement of discontinuation of **CentOS**. It is a downstream of **RHEL** that provides feature parity and binary compatibility with **RHEL**.
 - **Alma Linux** - A distribution created by the **CloudLinux** team after the announcement of discontinuation of **CentOS**. It is a downstream of **RHEL** that provides feature parity and binary compatibility with **RHEL**.
- **Arch Linux** - A community-driven distribution known for its simplicity and customizability. It is a *rolling release* distribution, which means that it is continuously updated. It is a bare-bones

1.1	Introduction	1
1.1.1	What is Linux?	1
1.1.2	Desktop Environments	2
1.1.3	Window Manager	3
1.1.4	Why Linux?	4
1.1.5	What is Shell?	4
1.1.6	Shell vs Terminal	4
1.1.7	Why the Command Line?	5
1.1.8	Command Prompt	6
1.1.9	Simple Commands in GNU Core Utils	6
1.2	Navigating the File System	21
1.2.1	What is a File System?	21
1.2.2	Why do we have proc and sys directories?	22
1.2.3	Paths	23
1.2.4	Basic Commands for Navigation	23
1.3	File Permissions	24
1.4	Inodes and Links	25
1.5	System Management and Information	26
1.6	Reading and Writing Files using cat	27
1.7	Types of Files	28
1.8	Types of Commands	28
1.9	File Manipulation	29

1: “Free software” means software that respects users’ freedom and community. Roughly, it means that the users have the freedom to run, copy, distribute, study, change and improve the software. Thus, “free software” is a matter of liberty, not price. To understand the concept, you should think of “free” as in “free speech,” not as in “free beer.” We sometimes call it “libre software,” borrowing the French or Spanish word for “free” as in freedom, to show we do not mean the software is gratis.

You may have paid money to get copies of a free program, or you may have obtained copies at no charge. But regardless of how you got your copies, you always have the freedom to copy and change the software, even to sell copies.

- GNU on Free Software




Figure 1.1: Linux Distributions Usage in 2024

2: GTK is a free software cross-platform widget toolkit for creating graphical user interfaces.

3: Ubuntu used to ship with Unity as the default desktop environment, but switched to GNOME in 2017.

4: Qt is cross-platform application development framework for creating graphical user interfaces.

distribution that lets the user decide which packages they want to install.

- **Manjaro** - A distribution based on Arch Linux which is known for its user-friendliness. It is a rolling release distribution that is easier to install for new users. It uses a different repository for packages with additional testing.
 - **EndeavourOS** - A distribution based on Arch Linux which is known for its simplicity and minimalism. It is a rolling release distribution that is easier to install for new users. It uses the same repository for packages as **Arch Linux**.
 - **Artix Linux** - It uses the **OpenRC** init system instead of **systemd**. It also offers other *init systems* like **runit**, **s6**, **dinit**.
- **openSUSE** - It is a free and open-source Linux distribution developed by the openSUSE project. It is offered in two main variations: **Tumbleweed**, an upstream rolling release distribution, and **Leap**, a stable release distribution which is sourced from SUSE Linux Enterprise.
- **Tumbleweed** - Rolling Release upstream.
 - **Leap** - Stable Release downstream.
- **Gentoo** - A distribution known for its customizability and performance. It is a source-based distribution, which means that the user compiles the software from source code. It is known for its performance optimizations for the user’s hardware.
- **Void** - It is an independent rolling-release Linux distribution that uses the X Binary Package System package manager, which was designed and implemented from scratch, and the **runit** init system. Excluding binary kernel blobs, a base install is composed entirely of free¹ software.

1.1.2 Desktop Environments

Definition 1.1.2 (Desktop Environment) A desktop environment is a collection of software designed to give functionality and a certain look and feel to a desktop operating system. It is a combination of a window manager, a file manager, a panel, and other software that provides a graphical user interface and utilities to a regular desktop user, such as volume and brightness control, multimedia applications, settings panels, etc. This is only required by desktop (and laptop) uses and are not present on server instances.

There are many desktop environments available for Linux, but the important ones are:

- **GNOME** - One of the most popular desktop environments for Linux. It is known for its simplicity and ease of use. It is the default desktop environment for many distributions, including Ubuntu. It is based on the **GTK** Toolkit.² Popular distros shipping by default with GNOME are Fedora, RHEL, CentOS, Debian, Zorin, and Ubuntu.³
- **KDE Plasma** - A highly customizable desktop environment based on the **Qt** Toolkit.⁴ Many distributions like Slackware and Open-

SUSE ship with KDE Plasma as the default desktop environment, and most others have the option to install with KDE Plasma. Ubuntu's KDE Plasma variant is called **Kubuntu**.

- ▶ **Xfce** - A lightweight desktop environment known for its speed and simplicity. It is based on the **GTK** Toolkit. It is used in many distributions like Xubuntu, Manjaro, and Fedora.
- ▶ **LXQt** - A lightweight desktop environment known for its speed and simplicity. It is based on the **Qt** Toolkit. It is used in many distributions like Lubuntu.
- ▶ **Cinnamon**
- ▶ **MATE**

It is important to note that although some distributions come pre-bundled with certain Desktop Environments, it doesn't mean that you cannot use another DE with it. DEs are simply packages installed on your distribution, and almost all the popular DEs can be installed on all distributions. Many distributions also come with multiple pre-bundled desktop environments due to user preferences. Most server distributions and some enthusiast distributions come with no pre-bundled desktop environment, and let the user determine which one is required, or if one is required.




Figure 1.2: Desktop Environment Usage in 2022

1.1.3 Window Manager

Definition 1.1.3 (Window Manager) A window manager is system software that controls the placement and appearance of windows within a windowing system in a graphical user interface. It is a part of the desktop environment, but can also be used standalone. It is responsible for the appearance and placement of windows, and can also provide additional functionality like virtual desktops, window decorations, window title bars, and tiling.

Although usually bundled with a desktop environment, many window managers are also standalone and installed separately by the user if they don't want to use all the application from a single desktop environment.

Some popular window managers are:

- ▶ **Openbox** - A lightweight window manager known for its speed and simplicity. It is used in many distributions like Lubuntu.
- ▶ **i3** - It is a tiling window manager⁵ which is usually one of the first window managers that users try when they want to move away from a desktop environment and to a tiling window manager.
- ▶ **awesome** - A tiling window manager that is highly configurable and extensible. It is written in Lua and is known for its beautiful configurations.
- ▶ **bspwm** - A tiling window manager. It is based on binary space partitioning.
- ▶ **dwm** - A dynamic tiling window manager that is known for its simplicity and minimalism. It is written in C and is highly configurable.

5: A tiling window manager is a window manager that automatically splits the screen into non-overlapping frames, which are used to display windows. Most desktop environments ship with a **floating** window manager instead, which users of other operating systems are more familiar with.

6: Although Linux is just a kernel and not an entire operating system, throughout this book I would be referring to **GNU/Linux**, the combination of **GNU core utilities** and the Linux kernel, as **Linux** in short.

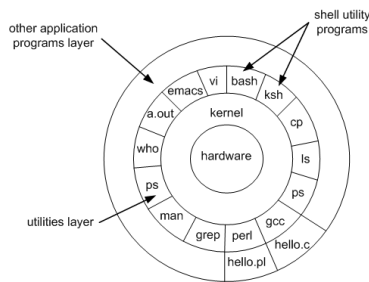


Figure 1.3: Operating System Onion Rings - The layers of an operating system

7: POSIX, or Portable Operating System Interface, is a set of standards that define the interfaces and environment that operating systems use to access POSIX-compliant applications. POSIX standards are based on the Unix operating system and were released in the late 1980s.

8: Fish is a non-POSIX compliant shell that is known for its features like auto-suggestion, syntax highlighting, and tab completions. Although a useful alternative to other shells for scripting, it should not be set as the default shell.

1.1.4 Why Linux?

You might be wondering “Why should I use Linux?” Most people use either **Windows** or **Mac** on their personal computers. Although these consumer operating systems get the job done, they don’t let the user completely control their own *hardware* and *software*. Linux⁶ is a free and open-source operating system that gives the user complete control over their system. It is highly customizable and can be tailored to the user’s needs. It is also known for its stability and security. It is used in almost all servers, supercomputers, and embedded systems. It is also used in many consumer devices like Android phones, smart TVs, and smartwatches.

In this course we will be covering how to navigate the linux file system, how to manage files, how to manage the system, and how to write scripts to automate tasks. In the later part of the course we go over concepts such as pattern matching and text processing.

This course does not go into details of the linux kernel, but rather attempts to make the reader familiar with the *GNU core utils* and able to navigate around a linux server easily.

1.1.5 What is Shell?

The **kernel** is the core of the operating system. It is responsible for managing the hardware and providing services to the user programs. The **shell** is the interface between the user and the kernel (Figure 1.3). Through the **shell** we can run many commands and utilities, as well as some inbuilt features of the shell.

Definition 1.1.4 (Shell) A shell is a command-line interpreter that provides a way for the user to interact with the operating system. It takes commands from the user and executes them. It is a program that provides the user with a way to interact with the operating system.

The most popular shell in Linux is the **bash** shell. It is the default shell in most distributions. It is a POSIX-compliant⁷ shell. There are many other shells available, such as **zsh**, **fish**⁸, **dash**, **csh**, **ksh**, and **tcsh**. Each shell has its own features and syntax, but most of the keywords and syntax are the same. In this course we will be covering only the **bash** shell and its syntax, but most of what we learn here is also applicable on other shells as well.

1.1.6 Shell vs Terminal

Definition 1.1.5 (Terminal) A terminal is a program that provides a way to interact with the shell. It is a program that provides a text-based interface to the shell. It is also known as a terminal emulator.

The terminal is the window that you see when you open a terminal program. It provides a way to interact with the shell. The shell is the program that interprets the commands that you type in the terminal.

The terminal is the window that you see, and the shell is the program that runs in that window. Whereas the shell is the application that is parsing your input and running the commands and keywords, the terminal is the application that lets you see the shell graphically. There are multiple different terminal emulators, providing a lot of customization and beautification to the terminal, as well as providing useful features such as *scroll back*, copying and pasting, and so on.

Some popular terminal emulators are:

- ▶ **gnome-terminal** - The default terminal emulator for the GNOME desktop environment.
- ▶ **konsole** - The default terminal emulator for the KDE desktop environment.
- ▶ **xfce4-terminal** - The default terminal emulator for the Xfce desktop environment.
- ▶ **alacritty** - A terminal emulator known for its speed and simplicity.
- ▶ **terminator** - A terminal emulator known for its features like splitting the terminal into multiple panes.
- ▶ **tilix** - A terminal emulator known for its features like splitting the terminal into multiple panes.
- ▶ **st** - A simple terminal emulator known for its simplicity and minimalism.
- ▶ **urxvt**
- ▶ **kitty**
- ▶ **terminology**

1.1.7 Why the Command Line?

Both the command line interface (CLI) and the graphical user interface (GUI) are simply shells over the operating system's kernel. They let you interact with the kernel, perform actions and run applications.

GUI:

The GUI requires a mouse and a keyboard, and is more intuitive and easier to use for beginners. But it is also slower and less efficient than the CLI. The GUI severely limits the user's ability to automate tasks and perform complex operations. The user can only perform those operations that the developers of the GUI have thought of and implemented.

CLI:

The CLI is faster and more efficient than the GUI as it lets the user use the keyboard to perform actions. Instead of clicking on pre-defined buttons, the CLI lets you construct your instruction to the computer using syntax and semantics. The CLI lets you combine simple commands that do one thing well to perform complex operations. The biggest advantage of the CLI is that it lets you automate tasks. It might be faster for some users to rename a file from **file1** to **file01** using the GUI, but it will always be faster to automate this using the CLI if you want to do this for thousands of files in the folder.

In this course we will be learning how to use the CLI to navigate the file system, manage files, manage the system, process text, and write scripts to automate tasks.

1.1.8 Command Prompt

The command prompt is the text that is displayed in the terminal to indicate that the shell is ready to accept commands. It usually ends with a \$ or a # symbol. The \$ symbol indicates that the shell is running as a normal user, and the # symbol indicates that the shell is running as the root user. The root user has complete control over the system and can perform any operation on the system.

An example of a command prompt is:

```
1 | username@hostname:~$
```

Here, username is the name of the user, hostname is the name of the computer, and \$ indicates that the shell is running as a normal user. The ~ symbol indicates that the current working directory is the user's home directory.⁹

9: The home directory is the directory where the user's files and settings are stored. It is usually located at /home/username. This can be shorted to ~ in the shell.

1.1.9 Simple Commands in GNU Core Utils




Figure 1.4: GNU Core Utils Logo

Definition 1.1.6 (GNU Core Utils) The GNU Core Utilities are the basic file, shell, and text manipulation utilities of the GNU operating system. These are the utilities that are used to interact with the operating system and perform basic operations.¹⁰

10: GNU Core Utils

The shell lets you simply type in the name of the command and press enter to run it. The shell then runs the command and displays the output. The shell also lets you pass arguments to the command to modify its behavior. Although the commands are simple, they are powerful and can be combined to perform complex operations.¹¹ Anything that can be done through a GUI can be done through the CLI, and often faster and more efficiently. Every GUI primitive has a CLI counterpart, such as copying, moving, renaming, deleting files, going inside a directory, etc.

Some basic commands in the core-utils are listed in Table Table 1.1 on the facing page.

11: The combination of commands to perform complex operations is called *pip-ing*. This will be covered later.

Command	Description
ls	List the contents of a directory
cd	Change the current working directory
pwd	Print the current working directory
mkdir	Create a new directory
rmdir	Remove a directory
touch	Create a new file
rm	Remove a file
cp	Copy a file
mv	Move a file
echo	Print a message
cat	Concatenate and display the contents of a file
less	Display the contents of a file one page at a time
head	Display the first few lines of a file
tail	Display the last few lines of a file
find	Find files and directories
locate	Find files and directories
which	Find the location of a command
uname	Print system information
ps	Display information about running processes
kill	Terminate a process
chmod	Change the permissions of a file
chown	Change the owner of a file
chgrp	Change the group of a file
date	Print the current date and time
cal, ncal	Print a calendar
df	Display disk space usage
du	Display disk usage
free	Display memory usage
top	Display system information
history	Display the command history
sleep	Pause the shell for a specified time
true	Do nothing, successfully
false	Do nothing, unsuccessfully
tee	Read from stdin and write to stdout and files
whoami	Print the current user
groups	Print the groups the user belongs to
clear	Clear the terminal screen
exit	Exit the shell

Table 1.1: Basic Commands in GNU Core Utils

pwd:

The `pwd` command prints the current working directory. The current working directory is the directory that the shell is currently in. The shell starts in the user's home directory when it is opened. The `pwd` command prints the full path of the current working directory.

```
1 | $ pwd
2 | /home/username
```

ls:

The `ls` command lists the contents of a directory. By default, it lists the contents of the current working directory. The `ls` command can take arguments to list the contents of a different directory.

```
1 | $ ls
2 | Desktop Documents Downloads Music Pictures Videos
```

12: Hidden files are files whose names start with a dot. They are hidden by default in the `ls` command.

We can also list hidden files¹² using the `-a` flag.

```
1 $ ls -a
2 .  .. .bashrc Desktop Documents Downloads Music Pictures
   Videos
```

Here, the `.` and `..` directories are special directories. The `.` directory is the current directory, and the `..` directory is the parent directory. The `.bashrc` file is a configuration file for the shell which is a hidden file.

`ls` can also list the details of the files using the `-l` flag.

```
1 $ ls -l
2 total 24
3 drwxr-xr-x 2 username group 4096 Mar  1 12:00 Desktop
4 drwxr-xr-x 2 username group 4096 Mar  1 12:00 Documents
5 drwxr-xr-x 2 username group 4096 Mar  1 12:00 Downloads
6 drwxr-xr-x 2 username group 4096 Mar  1 12:00 Music
7 drwxr-xr-x 2 username group 4096 Mar  1 12:00 Pictures
8 drwxr-xr-x 2 username group 4096 Mar  1 12:00 Videos
```




Figure 1.5: `ls -l` Output

As seen in Figure 1.5, the first column is the file type and permissions. The second column is the number of links to the file or directory. The third and fourth columns are the owner and group of the file or directory. The fifth column is the size of the file or directory. The sixth, seventh, and eighth columns are the last modified date and time of the file or directory. The ninth column is the name of the file or directory.¹³

13: More details about the file permissions and the file types will be covered later in the course.

14: An inode is a data structure on a filesystem on Linux and other Unix-like operating systems that stores all the information about a file except its name and its actual data. This includes the file type, the file's owner, the file's group, the file's permissions, the file's size, the file's last modified date and time, and the file's location on the disk. The inode is the reference pointer to the data in the disk.

We can also list the inode numbers¹⁴ using the `-i` flag.

```
1 $ ls -i
2 123456 Desktop 123457 Documents 123458 Downloads 123459 Music
   123460 Pictures 123461 Videos
```

Inodes will be discussed in detail later in the course.

cd:

The `cd` command changes the current working directory. It takes the path to the directory as an argument.

```
1 $ cd Documents
2 $ pwd
3 /home/username/Documents
```

The `cd` command can also take the `~` symbol as an argument to change to the user's home directory. This is the default behavior of the `cd` command when no arguments are passed.

```
1 $ cd
2 $ pwd
3 /home/username
```

15: This internally uses the `OLDPWD` environment variable to change the directory. More about variables will be covered later in the course.

If we want to go back to the previous directory, we can use the `-` symbol as an argument to the `cd` command.¹⁵

```
1 $ cd Documents
2 $ pwd
3 /home/username/Documents
4 $ cd -
5 $ pwd
6 /home/username
```

man:

How to remember so many flags and options for each of the commands?
The `man` command is used to display the manual pages for a command.

Definition 1.1.7 (Manual Pages) Manual pages are a type of software documentation that provides details about a command, utility, function, or file format. They are usually written in a simple and concise manner and provide information about the command's syntax, options, and usage.

```
1 | $ man ls
```

This will display the manual page for the `ls` command. The manual page is divided into sections, and you can navigate through the sections using the arrow keys. Press `q` to exit the manual page.

Example manual page:

```
1 |
2 | LS(1)
3 |
4 |                                     User Commands
5 |                                     LS(1)
6 |
7 | NAME
8 |     ls - list directory contents
9 |
10 | SYNOPSIS
11 |     ls [OPTION]... [FILE]...
12 |
13 | DESCRIPTION
14 |     List information about the FILES (the current directory by
15 |     default). Sort entries alphabetically if none of -cftuvSUX
16 |     nor --sort is specified.
17 |
18 |     Mandatory arguments to long options are mandatory for short
19 |     options too.
20 |
21 |     -a, --all
22 |         do not ignore entries starting with .
23 |
24 |     -A, --almost-all
25 |         do not list implied . and ..
26 |
27 | ...
```

The manual page provides information about the command, its syntax, options, and usage. It is a good practice to refer to the manual page of a command before using it.

To exit the manual page, press `q`.

There are multiple sections in the manual page, `man` takes the section number as an argument to display the manual page from that section.

```
1 | $ man 1 ls
```

This will display the manual page for the `ls` command from section 1. The details of the sections can be seen in Table 1.2 on the next page.

Table 1.2: Manual Page Sections

Section	Description
1	User Commands
2	System Calls
3	Library Functions
4	Special Files usually found in /dev
5	File Formats and conventions
6	Games
7	Miscellaneous
8	System Administration
9	Kernel Developer’s Manual

Man pages only provide information about the commands and utilities that are installed on the system. They do not provide information about the shell builtins or the shell syntax. For that, you can refer to the shell’s documentation or use the **help** command.

Some commands also have a `--help` flag that displays the usage and options of the command.

Some commands have their own **info** pages, which are more detailed than the **man** pages.

16: An useful video by Robert Elder about the differences between **man**, **info**, and **help** can be found on [YouTube](#).

To be proficient with shell commands, one needs to read the **man**, **info**, and **help** pages.¹⁶

Exercise 1.1.1 Run **man**, **info**, and `--help` on all the commands discussed in this section. Note the differences in the information provided by each of them. Read the documentations carefully and try to understand how each command works, and the pattern in which the documentations are written.

uname:

The **uname** command prints system information. It can take flags to print specific information about the system. By default, it prints only the kernel name.

17: Here **rex** is the hostname of the system, **6.8.2-arch2-1** is the kernel version, **x86_64** is the architecture, and **GNU/Linux** is the operating system.

```
1 | $ uname
2 | Linux
```

The `-a` flag prints all the system information.¹⁷

```
1 | $ uname -a
2 | Linux rex 6.8.2-arch2-1 #1 SMP PREEMPT_DYNAMIC Thu, 28 Mar 2024
   | 17:06:35 +0000 x86_64 GNU/Linux
```

ps:

The **ps** command displays information about running processes. By default, it displays information about the processes run by the current user that are running from a terminal.¹⁸

18: The **PID** is the process ID, the **TTY** is the terminal the process is running from, the **TIME** is the time the process has been running, and the **CMD** is the command that is running.

```
1 | $ ps
2 |  PID TTY          TIME CMD
3 | 12345 pts/0    00:00:00 bash
4 | 12346 pts/0    00:00:00 ps
```

There are a lot of flags that can be passed to the `ps` command to display more information about the processes. These will be covered in Chapter 3.

kill:

The `kill` command is used to terminate a process. It takes the process ID as an argument.

```
1 | $ kill 12345
```

The `kill` command can also take the **signal** number as an argument to send a signal to the process. For example, the `SIGKILL` signal can be sent to the process to terminate it.¹⁹

```
1 | $ kill -9 12345
```

19: The `SIGKILL` signal is used to terminate a process immediately. It cannot be caught or ignored by the process. It is numbered as 9.

mkdir:

The `mkdir` command is used to create new directories. It can take multiple directory names as arguments.

```
1 | $ mkdir a b c
2 | $ ls -F
3 | a/ b/ c/
```

Exercise 1.1.2 Run `man ls` to find out what the `-F` flag does, and why we used it in the above example.

touch:

The `touch` command is used to create new files. It can take multiple file names as arguments. If a file is already present, the `touch` command updates the last modified date and time of the file, but does not modify the contents of the file.

```
1 | $ touch file1 file2 file3
2 | $ ls -l
3 | -rw-r--r-- 1 username group 0 Mar  1 12:00 file1
4 | -rw-r--r-- 1 username group 0 Mar  1 12:00 file2
5 | -rw-r--r-- 1 username group 0 Mar  1 12:00 file3
6 | $ sleep 60
7 | $ touch file3
8 | $ ls -l
9 | -rw-r--r-- 1 username group 0 Mar  1 12:00 file1
10 | -rw-r--r-- 1 username group 0 Mar  1 12:00 file2
11 | -rw-r--r-- 1 username group 0 Mar  1 12:01 file3
```

Exercise 1.1.3 Notice the difference in the last modified date and time of the `file3` file from the other files. Also notice the `sleep` command used to pause the shell for 60 seconds.

rmdir:

The `rmdir` command is used to remove directories. It can take multiple directory names as arguments.

```

1 | $ mkdir a b c d
2 | $ rmdir a b c
3 | $ ls -F
4 | d/

```

Remark 1.1.1 The `rmdir` command can only remove empty directories. This is a safety feature so that users don't accidentally delete directories with files in them. To remove directories with files in them along with those files, use the `rm` command.

rm:

The `rm` command is used to remove files and directories. It can take multiple file and directory names as arguments.

```

1 | $ touch file1 file2 file3
2 | $ ls -F
3 | file1 file2 file3
4 | $ rm file1 file2
5 | $ ls -F
6 | file3

```

However, using `rm` to delete a directory will give an error.

```

1 | $ mkdir a
2 | $ rm a
3 | rm: cannot remove 'a': Is a directory

```

This is because the `rm` command does not remove directories by default. This is a safety feature to prevent users from accidentally deleting directories with files in them.

To remove directories along with their files, use the `-r` flag.

```

1 | $ rm -r a

```

To force the removal of files and directories without a confirmation, use the `-f` flag.

Warning 1.1.1 The `rm` command is a dangerous command. It does not move the files to the trash, but permanently deletes them. Be **extremely** careful when using the `rm` command. Only use the `-f` flag if you are absolutely sure that you want to delete the files.

To force `rm` to always ask for confirmation before deleting files, use the `-i` flag.

```

1 | $ rm -i file3
2 | rm: remove regular empty file 'file3'? y

```

cp:

The `cp` command is used to copy files. It takes the source file and the destination file as arguments.

```

1 | $ touch file1
2 | $ ls -F
3 | file1
4 | $ cp file1 file2
5 | $ ls -F

```

```
6 | file1 file2
```

The `cp` command can also take the `-r` flag to copy directories.

```
1 | $ mkdir a
2 | $ touch a/file1
3 | $ cp -r a b
4 | $ ls -R
5 | .:
6 | a/ b/
7 |
8 | ./a:
9 | file1
10 |
11 | ./b:
12 | file1
```

Exercise 1.1.4 Why did we use the `-R` flag in the above example? What does it do?

There are three ways to copy files using `cp`:

```
1 | SYNOPSIS
2 |     cp [OPTION]... [-T] SOURCE DEST
3 |     cp [OPTION]... SOURCE... DIRECTORY
4 |     cp [OPTION]... -t DIRECTORY SOURCE...
```

Exercise 1.1.5 There are three ways of running the `cp` command to copy a file. Here we have demonstrated only one. Read the manual page of the `cp` command to find out the other two ways and try them out yourself.

mv:

The `mv` command is used to move files. The syntax is similar to the `cp` command.²⁰ It is used to move files from one location to another, or to rename files.

20: This means that `mv` also has three ways of running it.

```
1 | $ touch file1
2 | $ ls -F
3 | file1
4 | $ mv file1 file2
5 | $ ls -F
6 | file2
```

Exercise 1.1.6 Create a directory `dir1` using the `mkdir` command, then create a file `file1` inside `dir1`. Now move (rename) the `dir1` directory to `dir2` using the `mv` command. The newly created directory should be named `dir2` and should contain the `file1` file. Were you require to use the `-r` flag with the `mv` command like you would have in `cp` command?

echo:

The `echo` command is used to print a message to the terminal. It can take multiple arguments and print them to the terminal.

```
1 | $ echo Hello, World!
2 | Hello, World!
```

The echo command can also take the -e flag to interpret backslash escapes.

```
1 | $ echo -e "Hello, \nWorld!"
2 | Hello,
3 | World!
```

Some escape characters in **echo** are listed in Table Table 1.3.

Table 1.3: Escape Characters in echo

Escape	Description
\\	backslash
\a	alert (BEL)
\b	backspace
\c	produce no further output
\e	escape
\f	form feed
\n	new line
\r	carriage return
\t	horizontal tab
\v	vertical tab
\0NNN	byte with octal value NNN (1 to 3 digits)
\xHH	byte with hexadecimal value HH (1 to 2 digits)

Exercise 1.1.7 Run the command `echo -e "\x41=\0101"` and try to understand the output and the escape characters used.

cat:

The cat command is used to concatenate and display the contents of files.

```
1 | $ cat file1
2 | Hello, World! from file1
```

The cat command can take multiple files as arguments and display their contents one after another.

```
1 | $ cat file1 file2
2 | Hello, World! from file1
3 | Hello, World! from file2
```

less:

Sometimes the contents of a file are too large to be displayed at once. Nowadays modern terminal emulators can scroll up and down to view the contents of the file, but actual **ttys** cannot do that. To view the contents of a file one page at a time, use the **less** command. **less** is a pager program that displays the contents of a file one page at a time.²¹

```
1 | $ less file1
```

To scroll up and down, use the arrow keys, or the j and k keys.²² Press q to exit the less command.

head:

21: **more** is another pager program that displays the contents of a file one page at a time. It is older than **less** and has fewer features. **less** is an improved version of **more** and is more commonly used. Due to this, it is colloquially said that "*less is more*", as it has more features.

22: Using j and k to move the cursor up and down is a common keybinding in many terminal applications. This originates from the *vi* text editor which will be covered later in the course.

The `head` command is used to display the first few lines of a file. By default, it displays the first 10 lines of a file.

```
1 $ head /etc/passwd
2 root:x:0:0:root:/root:/usr/bin/bash
3 bin:x:1:1:::/usr/bin/nologin
4 daemon:x:2:2:::/usr/bin/nologin
5 mail:x:8:12::/var/spool/mail:/usr/bin/nologin
6 ftp:x:14:11::/srv/ftp:/usr/bin/nologin
7 http:x:33:33::/srv/http:/usr/bin/nologin
8 nobody:x:65534:65534:Kernel Overflow User:/usr/bin/nologin
9 dbus:x:81:81:System Message Bus:/usr/bin/nologin
10 systemd-coredump:x:984:984:systemd Core Dumper:/usr/bin/
    nologin
11 systemd-network:x:982:982:systemd Network Management:/usr/bin/
    nologin
```

The `head` command can take the `-n` flag to display the first n lines of a file.²³

```
1 $ head -n 5 /etc/passwd
2 root:x:0:0:root:/root:/usr/bin/bash
3 bin:x:1:1:::/usr/bin/nologin
4 daemon:x:2:2:::/usr/bin/nologin
5 mail:x:8:12::/var/spool/mail:/usr/bin/nologin
6 ftp:x:14:11::/srv/ftp:/usr/bin/nologin
```

23: We can also directly run `head -5 /etc/passwd` to display the first 5 lines of the file.

Remark 1.1.2 Here we are listing the file `/etc/passwd` which contains information about the users on the system. The file will usually be present on all Unix-like systems and have a lot of system users.²⁴

tail:

The `tail` command is used to display the last few lines of a file. By default, it displays the last 10 lines of a file.

```
1 $ tail /etc/passwd
2 rtkit:x:133:133:RealtimeKit:/proc:/usr/bin/nologin
3 sddm:x:964:964:SDDM Greeter Account:/var/lib/sddm:/usr/bin/
    nologin
4 usbmux:x:140:140:usbmux user:/usr/bin/nologin
5 sayan:x:1000:1001:Sayan:/home/sayan:/bin/bash
6 qemu:x:962:962:QEMU user:/usr/bin/nologin
7 cups:x:209:209:cups helper user:/usr/bin/nologin
8 dhcpcd:x:959:959:dhcpcd privilege separation:/usr/bin/nologin
9 saned:x:957:957:SANE daemon user:/usr/bin/nologin
```

The `tail` command can take the `-n` flag to display the last n lines of a file.

```
1 $ tail -n 5 /etc/passwd
2 sayan:x:1000:1001:Sayan:/home/sayan:/bin/bash
3 qemu:x:962:962:QEMU user:/usr/bin/nologin
4 cups:x:209:209:cups helper user:/usr/bin/nologin
5 dhcpcd:x:959:959:dhcpcd privilege separation:/usr/bin/nologin
6 saned:x:957:957:SANE daemon user:/usr/bin/nologin
```

24: A system user is a user that is used by the system to run services and daemons. It does not belong to any human and usually is not logged into. System users have a user ID less than 1000.

Exercise 1.1.8 Notice that the UID (3rd column) of the `sayan` user is 1000. The last column is `/bin/bash` instead of `/usr/bin/nologin` like others. This is because it is a regular user and not a system user.

alias:

The `alias` command is used to create an alias for a command. An alias is a custom name given to a command that can be used to run the command.

25: Aliases are used to create shortcuts for long commands. They can also be used to create custom commands.

```
1 | $ alias ll='ls -l'
2 | $ ll
3 | total 24
4 | drwxr-xr-x 2 username group 4096 Mar  1 12:00 Desktop
5 | drwxr-xr-x 2 username group 4096 Mar  1 12:00 Documents
6 | drwxr-xr-x 2 username group 4096 Mar  1 12:00 Downloads
7 | drwxr-xr-x 2 username group 4096 Mar  1 12:00 Music
8 | drwxr-xr-x 2 username group 4096 Mar  1 12:00 Pictures
9 | drwxr-xr-x 2 username group 4096 Mar  1 12:00 Videos
```

The alias `ll` is created for the `ls -l` command.

Warning 1.1.2 Be careful when creating aliases. Do not create aliases for existing commands. This can lead to confusion and errors.

But this alias is temporary and will be lost when the shell is closed. To make the alias permanent, add it to the shell configuration file. For **bash**, this is the `.bashrc` file in the home directory.

```
1 | $ echo "alias ll='ls -l'" >> ~/.bashrc
```

This will add the alias to the `.bashrc` file. To make the changes take effect, either close the terminal and open a new one, or run the `source` command.

```
1 | $ source ~/.bashrc
```

Warning 1.1.3 Be careful when editing the shell configuration files. A small mistake can lead to the shell not working properly. Always keep a backup of the configuration files before editing them.

We can see the aliases that are currently set using the `alias` command.

```
1 | $ alias
2 | alias ll='ls -l'
```

We can also see what a particular alias expands to using the `alias` command with the alias name as an argument.

```
1 | $ alias ll
2 | alias ll='ls -l'
```

To remove an alias, use the `unalias` command.

```
1 | $ unalias ll
```

Exercise 1.1.9 Create an alias `la` for the `ls -a` command. Make it permanent by adding it to the `.bashrc` file. Check if the alias is set using the `alias` command.

Exercise 1.1.10 Create an alias `rm` for the `rm -i` command. Make it permanent by adding it to the `.bashrc` file. Check if the alias is set using the `alias` command. Try to delete a file using the `rm` command. What happens?

whoami:

The `whoami` command is used to print the username of the current user.

```
1 $ whoami
2 sayan
```

date:

The `date` command is used to print formatted date and time information. Without any arguments, it prints the current date and time.

```
1 $ date
2 Mon May 20 06:23:07 PM IST 2024
```

We can specify the date and time to be printed using the `-d` flag.

```
1 $ date -d "2020-05-20 00:30:45"
2 Wed May 20 12:30:45 AM IST 2020
3 $ date -d "2019-02-29"
4 date: invalid date '2019-02-29'
5 $ date -d "2020-02-29"
6 Sat Feb 29 12:00:00 AM IST 2020
```

Exercise 1.1.11 Why did we get an error when trying to print the date 2019-02-29?

We can also modify the format of the date and time using the `+` flag and different **format specifiers**. Some of the important format specifiers are listed in Table 1.4 on the next page. Rest of the format specifiers can be found in the date manual page.

```
1 $ date +"%Y-%m-%d %H:%M:%S"
2 2024-05-20 18:23:07
3 $ date +"%A, %B %d, %Y"
4 Monday, May 20, 2024
```

We can even mention relative dates and times using the `date` command.

```
1 $ date -d "next year"
2 Tue May 19 06:23:07 PM IST 2025
3 $ date -d "next month"
4 Thu Jun 20 06:23:07 PM IST 2024
5 $ date -d "tomorrow"
6 Tue May 21 06:23:07 PM IST 2024
7 $ date -d "yesterday"
8 Sun May 19 06:23:07 PM IST 2024
```

Table 1.4: Date Format Specifiers

Specifier	Description
%Y	Year
%m	Month
%d	Day
%H	Hour
%M	Minute
%S	Second
%A	Full weekday name
%B	Full month name
%a	Abbreviated weekday name
%b	Abbreviated month name

cal:

The `cal` command is used to print a calendar. By default, it prints the calendar of the current month.

```
1 | $ cal
2 |      May 2024
3 | Su Mo Tu We Th Fr Sa
4 |          1  2  3  4
5 |  5  6  7  8  9 10 11
6 | 12 13 14 15 16 17 18
7 | 19 20 21 22 23 24 25
8 | 26 27 28 29 30 31
```

We can specify the month and year to print the calendar of that month and year.

```
1 | $ cal 2 2024
2 |      February 2024
3 | Su Mo Tu We Th Fr Sa
4 |          1  2  3
5 |  4  5  6  7  8  9 10
6 | 11 12 13 14 15 16 17
7 | 18 19 20 21 22 23 24
8 | 25 26 27 28 29
```

There are multiple flags that can be passed to the `cal` command to display different types of calendars and of multiple months or of entire year.

Remark 1.1.3 In Linux, there are sometimes multiple implementations of the same command. For example, there are two implementations of the `cal` command, one in the **bsdmainutils** package, which is the **BSD** implementation and also includes another binary named **ncal** for printing the calendar in vertical format. The other implementation is in the **util-linux** package, which does not contain a **ncal** binary. The flags and the output of the `cal` command can differ between the two implementations.

free:

The `free` command is used to display the amount of free and used memory in the system.

```
1 | $ free
```

```

2      total      used      free      shared  buff/
   cache  available
3 Mem:    8167840  1234560    4567890    123456
      2367890    4567890
4 Swap:    2097148      0    2097148

```

The `free` command can take the `-h` flag to display the memory in human-readable format.

```

1 $ free -h
2      total      used      free      shared  buff/
   cache  available
3 Mem:    7.8Gi    1.2Gi    4.3Gi    120Mi
      2.3Gi    4.3Gi
4 Swap:    2.0Gi      0B    2.0Gi

```

df:

The `df` command is used to display the amount of disk space available on the filesystems.

```

1 $ df
2 Filesystem    1K-blocks    Used Available Use% Mounted on
3 /dev/sda1     12345678 1234567 11111111 10% /
4 /dev/sda2     12345678 1234567 11111111 10% /home

```

The `df` command can take the `-h` flag to display the disk space in human-readable format.

```

1 $ df -h
2 Filesystem    Size  Used Avail Use% Mounted on
3 /dev/sda1     12G  1.2G  9.9G  11% /
4 /dev/sda2     12G  1.2G  9.9G  11% /home

```

du:

The `du` command is used to display the disk usage of directories and files. By default, it displays the disk usage of the current directory.

```

1 $ du
2 4    ./Desktop
3 4    ./Documents
4 4    ./Downloads
5 4    ./Music
6 4    ./Pictures
7 4    ./Videos
8 28

```

The `du` command can take the `-h` flag to display the disk usage in human-readable format. The `-s` flag displays the total disk usage of the directory.

```

1 $ du -sh
2 28K  .

```

file:

The `file` command is used to determine the type of a file. It can take multiple file names as arguments.

```

1 $ file file1
2 file1: ASCII text
3 $ file /bin/bash

```

```

4 | /bin/bash: ELF 64-bit LSB shared object, x86-64, version 1 (SYSV
   | ), dynamically linked, interpreter /lib64/ld-linux-x86-64.so
   | .2, for GNU/Linux 3.2.0, BuildID[sha1]=1234567890abcdef,
   | stripped

```

groups:

The `groups` command is used to display the groups that the current user belongs to.

```

1 | $ groups
2 | sys wheel rfkill autologin sayan

```

These are some of the basic commands that are used in the terminal. Each of these commands has many more options and flags that can be used to customize their behavior. It is left as an exercise to the reader to explore the manual pages of these commands and try out the different options and flags.

In most terminal emulators, there are some basic shortcuts that can be used to make the terminal experience more efficient. Some of the basic shortcuts are listed in Table 1.5.

Table 1.5: Basic Shortcuts in Terminal

Shortcut	Description
Ctrl + C	Terminate the current process
Ctrl + D	Exit the shell
Ctrl + L	Clear the terminal screen
Ctrl + A	Move the cursor to the beginning of the line
Ctrl + E	Move the cursor to the end of the line
Ctrl + U	Delete from the cursor to the beginning of the line
Ctrl + K	Delete from the cursor to the end of the line
Ctrl + W	Delete the word before the cursor
Ctrl + Y	Paste the last deleted text
Ctrl + R	Search the command history
Ctrl + Z	Suspend the current process
Ctrl + \	Terminate the current process
Ctrl + S	Pause the terminal output
Ctrl + Q	Resume the terminal output

1.2 Navigating the File System

1.2.1 What is a File System?

Unlike Windows which has different drive letters for different partitions, Linux follows a unified file structure. The filesystem hierarchy is a tree of directories and files²⁶. The root²⁷ of the filesystem tree is the directory /. The basic filesystem hierarchy structure can be seen in Figure 1.6 and Table 1.6.

But what does so many directories mean? What do they do? What is the purpose of each directory?

Directory Path	Description
/	Root directory
/bin	Essential command binaries
/boot	Static files of the bootloader
/dev	Device files
/etc	Host-specific system configuration
/home	User home directories
/lib	Essential shared libraries and kernel modules
/media	Mount point for removable media
/mnt	Mount point for mounting a filesystem temporarily
/opt	Add-on application software packages
/proc	Virtual filesystem providing process information
/root	Home directory for the root user
/run	Run-time variable data
/sbin	Essential system binaries
/srv	Data for services provided by the system
/sys	Kernel and system information
/tmp	Temporary files
/usr	Secondary hierarchy
/var	Variable data

Some directories do not store data on the disk, but are used to store information about the system. These directories are called *virtual* directories. For example, the **/proc** directory is a virtual directory that provides information about the running processes. The **/sys** directory is another virtual directory that provides information about the system. The **/tmp** is a *volatile* directory whose data is deleted as soon as the system is turned off. The **/run** directory is another volatile directory that stores runtime data.

Rest directories are stored on the disk. The reason for having so many directories is to categorize the type of files they store. For example, all the executable binaries of different applications and utilities installed in the system is stored in **/bin** and **/sbin** directories. All the shared libraries installed on the system are stored in **/lib** directory. Sometimes some applications are installed in **/opt** directory which are not installed directly by the package manager.²⁸

We also need to store the user's documents and files. This is done in the **/home** directory. Each user has their own directory in the **/home** directory. The root user's directory is **/root**. All the application's configuration files are stored in the user's home directory in the **/home** directory itself. This

26: A non-directory is a leaf node of a tree.

27: The root of a tree is the first node from which the tree originates. A tree can have only one root.

Table 1.6: Linux Filesystem Hierarchy




Figure 1.6: Linux Filesystem Hierarchy

28: In Linux, you do not install applications by downloading them from the internet and running an installer like in Windows. You use a package manager to install applications. The package manager downloads the application from the internet and installs it on your system automatically. This way the package manager can also keep track of the installed applications and their dependencies and whether they should be updated. This is similar to the *Play Store* on mobile phones.

separation of application binary and per-user application settings helps people to easily change systems but keep their own **/home** directory constant and in turn, also all their application settings.

Some settings however needs to be applied system-wide and for all users. These settings are stored in the **/etc** directory. This directory contains all the system-wide configuration files.

To boot up the system, the bootloader needs some files. These files are stored in the **/boot** directory.²⁹ The bootloader is the first program that runs when the computer is turned on. It loads the operating system into memory and starts it.

Although the file system is a unified tree hierarchy, this doesn't mean that we cannot have multiple partitions on Linux: au contraire, it is easier to manage partitions on Unix. We simply need to mention which empty directory in the hierarchy should be used to mount a partition. As soon as that partition is mounted, it gets populated with the data stored on that disk with all the files and subdirectories, and when the device is unmounted the directory again becomes empty. Although a partition can be mounted on any directory, there are some dedicated folders in **/** as well for this purpose. For example, the **/mnt** directory is used to mount a filesystem temporarily, and the **/media** directory is used to mount removable media like USB drives, however it is not required to strictly follow this.

Finally, the biggest revelation in Linux is that, everything is a file. Not only are all the system configurations stored as **plain text** files which can be read by humans, but the processes running on your system are also stored as files in **proc**. Your kernel's interfaces to the applications or users are also simple files stored in **sys**. Biggest of all, even your hardware devices are stored as files in **dev**.³⁰

The **/usr** directory is a secondary hierarchy that contains subdirectories similar to those present in **/**. This was created as the olden system had started running out of disk space for the **/bin** and **/lib** directories. Thus another directory named **usr** was made, and subdirectories like **/usr/bin** and **/usr/lib** were made to store half of the binaries and libraries. There wasn't however any rigid definition of which binary should go where. Modern day systems have more than enough disk space to store everything on one partition, thus the **/bin** and **/lib** don't really exist any more. If they do, they are simply shortcuts³¹ to the **/usr/bin** and **/usr/lib** directories which are still kept for *backwards compatibility*.

These can also be loosely classified into *sharable* and *non-sharable* directories and *static* and *variable* directories as shown in Table 1.7.

Table 1.7: Linux Filesystem Directory Classification

	Sharable	Non-sharable
Static	/usr , /opt	/etc , /boot
Variable	/var/spool	/tmp , /var/log

1.2.2 Why do we have **proc** and **sys** directories?

The **/proc** tree originated in System V Unix³², where it only gave information about each running process, using a **/proc/\$PID/** format.

29: Modern systems use **UEFI** instead of **BIOS** to boot up the system. The bootloader is stored in the **/boot/EFI** directory or in the **/efi** directory directly.

30: Device files are not stored as normal files on the disk, but are special files that the kernel uses to communicate with the hardware devices. These are either **block** or **character** devices. They are used to read and write data to and from the hardware devices.

31: Shortcuts in Linux are called *symbolic links* or *symlinks*.

32: Unix System V is one of the first commercial versions of the Unix operating system. It was originally developed by AT&T and first released in 1983.

Linux greatly extended that, adding all sorts of information about the running kernel's status. In addition to these read-only information files, Linux's `/proc` also has writable virtual files that can change the state of the running kernel. BSD³³ type OSes generally do not have `/proc` at all, so much of what you find under `proc` is non-portable.

The intended solution for this mess in Linux's `/proc` is `/sys`. Ideally, all the non-process information that got crammed into the `/proc` tree should have moved to `/sys` by now, but historical inertia has kept a lot of stuff in `/proc`. Often there are two ways to effect a change in the running kernel: the old `/proc` way, kept for backwards compatibility, and the new `/sys` way that you're supposed to be using now.

33: BSD, or Berkeley Software Distribution, is a Unix-like operating system that was developed at the University of California, Berkeley. It was first released in 1977 and was based on the original Unix source code from AT&T. BSD is not linux, it is a totally different kernel, with similar core utils to GNU.

1.2.3 Paths

Whenever we open a terminal on a Linux system, we are placed in a directory. This is called the *current working directory*. All shells and applications have a current working directory from where they are launched.

To refer to and identify the directory you are talking about, we use a **path**.

Definition 1.2.1 (Path) Path is a traversal in the filesystem tree. It is a way to refer to a file or directory.

Absolute Path:

The traversal to the directory from the root directory is called the **absolute path**. For example, if we want to refer to the directory named **alex** inside the directory **home** in the root of the file system, then it is qualified as:

```
1 | /home/alex
```

Relative Path:

The traversal to the directory from the current working directory is called the **relative path**. For example, if we want to refer to the directory named **alex** inside the directory **home** from the `/usr/share` directory, then it will be qualified as:

```
1 | ../../home/alex
```

Remark 1.2.1 The `..` in the path refers to the parent directory. It is used in relative paths to refer to directories whose path requires travelling up the tree.




Figure 1.7: Relative Path

1.2.4 Basic Commands for Navigation

The file system can be navigated in the Linux command line using the following commands:

- **pwd**: Print the current working directory
- **ls**: List the contents of the current directory

34: **pushd** and **popd** are useful for quickly switching between directories in scripts.

- ▶ **cd**: Change the current working directory
- ▶ **mkdir**: Create a new directory
- ▶ **rmdir**: Remove a directory
- ▶ **touch**: Create a new file
- ▶ **rm**: Remove a file
- ▶ **pushd**: Push the current directory to a stack
- ▶ **popd**: Pop the current directory from a stack³⁴

More details about these commands can be found in their respective man pages. For example, to find more about the **ls** command, you can type `man ls`.

Question 1.2.1 What is the command to list the contents of the current directory?

Answer 1.2.1 `ls`

Question 1.2.2 What is the command to list the contents of the current directory including hidden files?

Answer 1.2.2 `ls -a`

Question 1.2.3 What is the command to list the contents of the current directory in a long list format? (show permissions, owner, group, size, and time)

Answer 1.2.3 `ls -l`

Question 1.2.4 What is the command to list the contents of the current directory in a long list format and also show hidden files?

Answer 1.2.4 `ls -al` or `ls -la` or `ls -l -a` or `ls -a -l`

Question 1.2.5 The output of `ls` gives multiple files and directories in a single line. How can you make it print one file or directory per line?

Answer 1.2.5 `ls -1`

This can also be done by passing the output of `ls` to `cat` or storing the output of `ls` in a file and then using `cat` to print it. We will see these in later weeks.³⁵

35: that is a one, not an L

1.3 File Permissions

Question 1.3.1 How to list the permissions of a file?

Answer 1.3.1 `ls -l`

The permissions are the first 10 characters of the output.

`stat -c %A filename` will list only the permissions of a file.

There are other format specifiers of `stat` to show different statistics which can be found in `man stat`.

Question 1.3.2 How to change permissions of a file? Let's say we want to change `file1`'s permissions to `rwxr-xr-`. What is the octal form of that?

Answer 1.3.2 `chmod u=rwx,g=rx,o=r file1` will change the permissions of `file1`

The octal form of `rwxr-xr-` is 754.

So we can also use `chmod 754 file1`

Providing the octal is same as using `=` to set the permissions.

We can also use `+` to add permissions and `-` to remove permissions.

1.4 Inodes and Links

Question 1.4.1 How to list the inodes of a file?

Answer 1.4.1 `ls -li` will list the inodes of a file. The inodes are the first column of the output of `ls -li`. This can be combined with other flags like `-l` or `-a` to show more details.

Question 1.4.2 How to create soft link of a file?

Answer 1.4.2 `ln -s sourcefile targetfile` will create a soft link of `sourcefile` named `targetfile`. The soft link is a pointer to the original file.

Definition 1.4.1 (Soft Links) Soft Links are special kinds of files that just store the path given to them. Thus the path given while making soft links should either be an absolute path, or relative **from** the location of the soft link **to** the location of the original file. It should not be relative from current working directory.^a

^a This is a common mistake.

Question 1.4.3 How to create hard link of a file?

Answer 1.4.3 `ln sourcefile targetfile` will create a hard link of `sourcefile` named `targetfile`. The hard link is same as the original file. It does not depend on the original file anymore after creation. They are equals, both are hardlinks of each other. There is no parent-child relationship. The other file can be deleted and the original file will still work.

Definition 1.4.2 (Hard Links) Hard Links are just pointers to the same inode. They are the same file. They are not pointers to the path of the file. They are pointers to the file itself. They are not affected by the deletion of the other file. When creating a hard link, you need to provide the path of the original file, and thus it has to be either absolute path, or relative from the current working directory, not relative from the location of the hard link.

Question 1.4.4 How to get the real path of a file?
Assume three files:

- ▶ **file1** is a soft link to **file2**
- ▶ **file2** is a soft link to **file3**
- ▶ **file3** is a regular file

Real path of all these three should be the same. How to get that?

Answer 1.4.4 `realpath filename` will give the real path of `filename`. You can also use `readlink -f filename` to get the real path.

1.5 System Management and Information

Question 1.5.1 How to print the current date and time in some custom format?

Answer 1.5.1 `date -d today +%Y-%m-%d` will print the current date in the format `YYYY-MM-DD`. The format can be changed by changing the format specifiers. The format specifiers are given in the `man date` page. The `-d today` can be dropped, but is mentioned to show that the date can be changed to any date. It can be strings like `'2024-01-01'` or `'5 days ago'` or `'yesterday'`, etc.

Question 1.5.2 How to print the kernel version of your system?

Answer 1.5.2 `uname -r` will print the kernel version of your system. `uname` is a command to print system information. The `-r` flag is to print the kernel release. There are other flags to print other system information.

We can also run `uname -a` to get all fields and extract only the kernel

info using commands taught in later weeks.

Question 1.5.3 How to see how long your system is running for? What about the time it was booted up?

Answer 1.5.3 `uptime` will show how long the system is running for.
`uptime -s` will show the time the system was booted up.
 The `-s` flag is to show the time of last boot.

Question 1.5.4 How to see the amount of free memory? What about free hard disk space? If we are unable to understand the big numbers, how to convert them to human readable format? What is difference between MB and MiB?

Answer 1.5.4 `free` will show the amount of free memory.
`df` will show the amount of free hard disk space.
`df -h` and `free -h` will convert the numbers to human readable format.
 MB is Megabyte, and MiB is Mebibyte.
 1 MB = 1000 KB, 1 GB = 1000 MB, 1 TB = 1000 GB, this is SI unit.
 1 MiB = 1024 KiB, 1 GiB = 1024 MiB, 1 TiB = 1024 GiB, this is 2^{10} unit.

1.6 Reading and Writing Files using cat

Question 1.6.1 Can we print contents of multiple files using a single command?

Answer 1.6.1 `cat file1 file2 file3` will print the contents of `file1`, `file2`, and `file3` in the order given. The contents of the files will be printed one after the other.

Question 1.6.2 Can `cat` also be used to write to a file?

Answer 1.6.2 Yes, `cat > file1` will write to `file1`. The input will be taken from the terminal and written to `file1`. The input will be written to `file1` until the user presses `Ctrl+D` to indicate end of input. This is redirection, which we see in later weeks.

Question 1.6.3 `ls` can only show files and directories in the **cwd**³⁶, not subdirectories. True or False?

^{36:} **cwd** means **C**urrent **W**orking **D**irectory

Answer 1.6.3 False. `ls` can show files and directories in the `cwd`, and also in subdirectories. The `-R` flag can be used to show files and

directories in subdirectories, recursively.

1.7 Types of Files

Question 1.7.1 What types of files are possible in a linux file system?

Answer 1.7.1 There are 7 types of files in a linux file system:

- ▶ Regular Files (starts with -)
- ▶ Directories (starts with d)
- ▶ Symbolic Links (starts with l)
- ▶ Character Devices (starts with c)
- ▶ Block Devices (starts with b)
- ▶ Named Pipes (starts with p)
- ▶ Sockets (starts with s)

Question 1.7.2 How to know what kind of file a file is? Can we determine using its extension? Can we determine using its contents? What does *MIME* mean? How to get that?

Answer 1.7.2 The file command can be used to determine the type of a file.

The extension of a file does not determine its type.

The contents of a file can be used to determine its type.

MIME stands for Multipurpose Internet Mail Extensions.

It is a standard that indicates the nature and format of a document.

`file -i filename` will give the MIME type of filename.

1.8 Types of Commands

Question 1.8.1 How to create aliases? How to make them permanent? How to unset them?

Answer 1.8.1 `alias name='command'` will create an alias.

`unalias name` will unset the alias.

To make them permanent, add the alias to the `~/.bashrc` file.

The `~/.bashrc` file is a script that is executed whenever a new terminal is opened.

Question 1.8.2 How to run the normal version of a command if it is aliased?

Answer 1.8.2 `\command` will run the normal version of command if it is aliased.

Question 1.8.3 What is the difference between `which`, `what is`, `where is`, `locate`, and `type`?

Answer 1.8.3 Each of the commands serve a different purpose:

- ▶ `which` will show the path of the command that will be executed.
- ▶ `what is` will show a short description of the command.
- ▶ `where is` will show the location of the command, source files, and man pages.
- ▶ `locate` is used to find files by name.
- ▶ `type` will show how the command will be interpreted by the shell.

Exercise 1.8.1 Find the path of the `true` command using `which`. Find a short description of the `true` command using `what is`. Is the executable you found actually the one that is executed when you run `true`? Check using `type true`

Definition 1.8.1 (Types of commands) A command can be an alias, a shell built-in, a shell function, keyword, or an executable. The `type` command will show which type the command is.

- ▶ **alias:** A command that is an alias to another command defined by the user or the system.
- ▶ **builtin:** A shell built-in command is a command that is built into the shell itself. It is executed internally by the shell.
- ▶ **file:** An executable file that is stored in the file system. It has to be stored somewhere in the `PATH` variable.
- ▶ **function:** A shell function is a set of commands that are executed when the function is called.
- ▶ **keyword:** A keyword is a reserved word that is part of the shell syntax. It is not a command, but a part of the shell syntax.

1.9 File Manipulation

Question 1.9.1 How to list only first 10 lines of a file? How about first 5? Last 5? How about lines 105 to lines 152?

Answer 1.9.1 `head filename` will list the first 10 lines of `filename`.
`head -n 5 filename` will list the first 5 lines of `filename`.
`tail -n 5 filename` will list the last 5 lines of `filename`.
`head -n 152 filename | tail -n 48` will list lines 105 to 152 of `filename`. This uses `|` which is a pipe, which we will see in later weeks.

Question 1.9.2 Do you know how many lines a file contains? How can we count it? What about words? Characters?

Answer 1.9.2 `wc filename` will count the number of lines, words, and characters in `filename`.

`wc -l filename` will count the number of lines in `filename`.

`wc -w filename` will count the number of words in `filename`.

`wc -c filename` will count the number of characters in `filename`.

Question 1.9.3 How to delete an empty directory? What about non-empty directory?

Answer 1.9.3 `rmdir dirname` will delete an empty directory.

`rm -r dirname` will delete a non-empty directory.

Question 1.9.4 How to copy an entire folder to another name? What about moving?

Why the difference in flags?

Answer 1.9.4 `cp -r sourcefolder targetfolder` will copy an entire folder to another name.

`mv sourcefolder targetfolder` will move an entire folder to another name.

The difference in flags is because `cp` is used to copy, and `mv` is used to move or rename a file or folder. The `-r` flag is to copy recursively, and is not needed for `mv` as it is not recursive and simply changes the name of the folder (or the path).

2.1 Creating Variables

Definition 2.1.1 (Environment Variables) An environment variable is a variable that is accessible to all processes running in the environment. It is a key-value pair. It is created using the `export` command. It can be accessed using the `$` followed by the name of the variable (e.g. `$HOME`) or using the `printenv` command. They are part of the environment in which a process runs. For example, a running process can query the value of the `TEMP` environment variable to discover a suitable location to store temporary files, or the `HOME` or `USER` variable to find the directory structure owned by the user running the process.

Definition 2.1.2 (Shell Variables) A shell variable is a variable that is only accessible to the shell in which it is created. It is a key-value pair. It is created using the `=` operator. It can be accessed using the `$` followed by the name of the variable (e.g. `$var`). They are local to the shell in which they are created.

Question 2.1.1 How to create a shell variable? How to create an environment variable?

Answer 2.1.1 `var=value` will create a shell variable.
`export var=value` will create an environment variable.

Question 2.1.2 How to set the prompt string? How to set the secondary prompt string?

Answer 2.1.2 `PS1="newprompt"` will set the prompt string.
`PS2="newprompt"` will set the secondary prompt string. This is useful when writing multi-line commands.

Question 2.1.3 Where are commands searched for? How to add a new path to search for commands? What symbol is used to separate the paths?

Answer 2.1.3 The `PATH` variable stores the paths to search for commands.
`export PATH=$PATH:newpath` will add `newpath` to the `PATH` variable.

2.1 Creating Variables	31
2.2 Printing Variables to the Terminal	32
2.3 Bash Flags	33
2.4 Signals	33
2.5 Variable Manipulation . .	34
2.6 Arrays	36

The colon : is used to separate the paths.

2.2 Printing Variables to the Terminal

Question 2.2.1 List the frequently used shell variables already existing in your shell.

Answer 2.2.1 Frequently used shell variables:

- ▶ **USER** stores the currently logged in user.
- ▶ **HOME** stores the home directory of the user.
- ▶ **PWD** stores the current working directory.
- ▶ **SHELL** stores the path of the shell being used.
- ▶ **PATH** stores the paths to search for commands.
- ▶ **PS1** stores the prompt string for the shell.
- ▶ **PS2** stores the secondary prompt string for the shell
- ▶ **HOSTNAME** stores the network name of the system

Question 2.2.2 List the special shell variables useful in scripts.

Answer 2.2.2 Special shell variables useful in scripts:

- ▶ **\$0** stores the name of the script.
- ▶ **\$1, \$2, \$3, ...** store the arguments to the script.
- ▶ **\$#** stores the number of arguments to the script.
- ▶ **\$*** stores all the arguments to the script as a single string.
- ▶ **@** stores all the arguments to the script as array of strings.
- ▶ **?** stores the exit status of the last command.
- ▶ **\$\$** stores the process id of the current shell.
- ▶ **!** stores the process id of the last background command.
- ▶ **-\$** stores the current options set for the shell.
- ▶ **IFS** stores the Internal Field Separator.
- ▶ **\$LINENO** stores the current line number of the script.
- ▶ **\$RANDOM** stores a random number.
- ▶ **\$SECONDS** stores the number of seconds the script has been running.

Question 2.2.3 How to print things to the terminal? How to print the value of a variable? How to print escape characters? How to suppress the newline?

Answer 2.2.3 `echo "Hello World"` will print Hello World to the terminal.

`echo $var` will print the value of var to the terminal.

`echo -e "Hello\nWorld"` will print Hello and World in two lines.

`echo -n "Hello"` will print Hello without a newline.

Question 2.2.4 How to print all the environment variables? Only some variables?

Answer 2.2.4 `printenv` will print all the environment variables.
`env` will also print all the environment variables.
`printenv var1 var2 var3` will print only `var1`, `var2`, and `var3`.

Question 2.2.5 How to escape a character in a string? How to escape a variable in a string?

Answer 2.2.5 `echo -e "Hello\nWorld"` will print Hello and World in two lines.
`echo "Hello $var"` will print Hello and the value of `var`.

2.3 Bash Flags

Question 2.3.1 How to list the bash flags set in the shell? How to set or unset them? What does the default flags mean? (himBHs)

Answer 2.3.1 `echo $-` will list the bash flags set in the shell.
`set -o flag` will set the flag.
`set +o flag` will unset the flag.
The default flags are:
h: locate and remember (hash) commands as they are looked up.
i: **interactive** shell
m: monitor the jobs and report changes
B: **braceexpand** - expand the expression in braces
H: **histexpand** - expand the history command
s: Read commands from the standard input.

2.4 Signals

Remark 2.4.1 If you press `Ctrl+S` in the terminal, some terminals might stop responding. You can resume it by pressing `Ctrl+Q`. This is because `ixoff` is set. You can unset it using `stty -ixon`. This is a common problem with some terminals, thus it can be placed inside the `/.bashrc` file.

Other Ctrl+Key combinations:

- ▶ **Ctrl+C**: Interrupt the current process, this sends the `SIGINT` signal.
- ▶ **Ctrl+D**: End of input, this sends the `EOF` signal.
- ▶ **Ctrl+L**: Clear the terminal screen.
- ▶ **Ctrl+Z**: Suspend the current process, this sends the `SIGTSTP` signal.

- ▶ **Ctrl+R**: Search the history of commands using reverse-i-search.
- ▶ **Ctrl+T**: Swap the last two characters
- ▶ **Ctrl+U**: Cut the line before the cursor
- ▶ **Ctrl+V**: Insert the next character literally
- ▶ **Ctrl+W**: Cut the word before the cursor
- ▶ **Ctrl+Y**: Paste the last cut text

2.5 Variable Manipulation

Question 2.5.1 How to remove a variable that is set? How to check if a variable is present?

Answer 2.5.1 `unset var` will remove the variable `var`.
`echo $var` will print if the variable `var` is present.
`test -v var` will check if the variable `var` is present.
`test -z $var+x` will check if the variable `var` is present.

Question 2.5.2 How to print a default value if variable is absent? if present? how to set a default value if variable is absent?

Answer 2.5.2 `echo $var:-default` will print default if `var` is absent.
`echo $var:+default` will print default if `var` is present.
`echo $var:=default` will set `var` to default if `var` is absent.

Remark 2.5.1 `echo $var:-default` will print default if `var` is absent **or** empty.
`echo $var-default` will print default if `var` is absent, but not if `var` is empty.
`echo $var:+default` will print default if `var` is present **and** not empty.
`echo $var+default` will print default if `var` is present, regardless of whether `var` is empty or not.

Question 2.5.3 How to print the number of characters in a variable? How about a substring of the variable?

Answer 2.5.3 `echo $#var` will print the number of characters in `var`.
`echo $var:0:5` will print the first 5 characters of `var`.

Remark 2.5.2 The second number in the substring is the length of the substring. Thus, `echo $var:5:3` will print the 3 characters starting

from the 5th character. If the length of substring is more than the length of the variable, it will print till the end. Thus, `echo $var:5:100` will print the characters from the 5th character till the end.

Definition 2.5.1 (Variable Match and Delete) We can match some part of the variable and delete it from left or right:

`${var%pattern}` will delete the shortest match of `pattern` from the end of `var`.

`${var%%pattern}` will delete the longest match of `pattern` from the end of `var`.

`${var#pattern}` will delete the shortest match of `pattern` from the start of `var`.

`${var##pattern}` will delete the longest match of `pattern` from the start of `var`.

Question 2.5.4 If we have a variable with value "abc.def.ghi.xyz". How can we get the following:

- ▶ "abc.def.ghi"
- ▶ "abc"
- ▶ "def.ghi.xyz"
- ▶ "xyz"

Answer 2.5.4 The variable is `var="abc.def.ghi.xyz"`.

- ▶ `echo ${var%.*}` will print "abc.def.ghi".
- ▶ `echo ${var%%.*}` will print "abc".
- ▶ `echo ${var#*.}` will print "def.ghi.xyz".
- ▶ `echo ${var##*.}` will print "xyz".

Question 2.5.5 How to replace some part of a variable by something else? How to do for all occurrences?

Answer 2.5.5 `echo ${var/pattern/string}` will replace the first occurrence of `pattern` with `string`.

`echo ${var//pattern/string}` will replace all occurrences of `pattern` with `string`.

Question 2.5.6 How to make a variable lowercase? How to make it uppercase? How to do for entire string? How to do for first character? How to make a variable Sentence case? (First letter is capital, rest lower)

Answer 2.5.6 The following commands can be used to manipulate the case of a variable:

```
1 echo ${var,,} # will make the variable lowercase.
2 echo ${var^^} # will make the variable uppercase.
3 echo ${var,} # will make the first character lowercase.
```

```

4 echo ${var^} # will make the first character uppercase.
5 lower=${var,,} # will store the lowercase variable in lower.
6 echo ${lower^} # will make the first character uppercase. (
  Sentence case)

```

Question 2.5.7 How to make a variable integer only? How to make it lowercase only? Uppercase only? Read only?

Answer 2.5.7 declare -i var will make the variable integer only.
 declare -l var will make the variable lowercase only.
 declare -u var will make the variable uppercase only.
 declare -r var will make the variable read only.

Question 2.5.8 How to remove those restrictions from a variable?

Answer 2.5.8 declare +i var will remove the integer only restriction.
 declare +l var will remove the lowercase only restriction.
 declare +u var will remove the uppercase only restriction.
 declare +r var will not work, as read only cannot be removed.

Question 2.5.9 How to store output of a command in a variable?

Answer 2.5.9 var=\$(command) will store the output of command in var.

2.6 Arrays

Question 2.6.1 How to create an array? How to access an element of an array? How to print all elements of an array?

Answer 2.6.1 arr=(1 2 3 4 5) will create an array.
 echo \${arr[2]} will access the 3rd element of the array.
 echo \${arr[@]} will print all elements of the array.

Question 2.6.2 How to get the number of elements in an array? How to get the indices of the array? How to add an element to the array?

Answer 2.6.2 echo \${#arr[@]} will print the number of elements in the array.
 echo \${!arr[@]} will print the indices of the array.
 arr+=(6) will add the element 6 to the array.

Question 2.6.3 What are associative arrays? How to declare and use them?

Answer 2.6.3 Associative arrays are arrays with key-value pairs.

`declare -A arr` will declare an associative array.

`arr[key]=value` will set the value of key to value.

`echo ${arr[key]}` will print the value of key.

Question 2.6.4 How to store output of a command in an array?

Answer 2.6.4 `arr=($(command))` will store the output of command in `arr`.

3.1 What are processes?

3.1 What are processes? 39

Definition 3.1.1 (Process) A process is an instance of a program that is being executed. It contains the program code and its current activity. Depending on the operating system (OS), a process may be made up of multiple threads of execution that execute instructions concurrently. Several processes may be associated with the same program; for example, opening up several instances of the same program often means more than one process is being executed. Each process has its own 'process id' or **PID** to uniquely identify it.

Question 3.1.1 How to get a snapshot of all the processes running? What are the commonly used flags used with it?

Answer 3.1.1 The **ps** command is used to get a snapshot of all the processes running.

- ▶ **ps** will get a snapshot of all the processes running.
- ▶ **ps -e** will show all the processes.
- ▶ **ps -f** will show full format listing.
- ▶ **ps -l** will show long format listing.
- ▶ **ps -u** will show user-oriented format listing.
- ▶ **ps -x** will show processes without controlling terminals.
- ▶ **ps -a** will show all processes with a terminal.
- ▶ **ps -A** will show all processes
- ▶ **ps aux** is a common command to see all processes
- ▶ **ps -forest** will show the processes in a tree form

APPENDIX

Notation

The next list describes several symbols that will be later used within the body of the document.

c Speed of light in a vacuum inertial frame

h Planck constant

Greek Letters with Pronunciations

Character	Name	Character	Name
α	alpha <i>AL-fuh</i>	ν	nu <i>NEW</i>
β	beta <i>BAY-tuh</i>	ξ, Ξ	xi <i>KSIGH</i>
γ, Γ	gamma <i>GAM-muh</i>	\omicron	omicron <i>OM-uh-CRON</i>
δ, Δ	delta <i>DEL-tuh</i>	π, Π	pi <i>PIE</i>
ϵ	epsilon <i>EP-suh-lon</i>	ρ	rho <i>ROW</i>
ζ	zeta <i>ZAY-tuh</i>	σ, Σ	sigma <i>SIG-muh</i>
η	eta <i>AY-tuh</i>	τ	tau <i>TOW (as in cow)</i>
θ, Θ	theta <i>THAY-tuh</i>	υ, Υ	upsilon <i>OOP-suh-LON</i>
ι	iota <i>eye-OH-tuh</i>	ϕ, Φ	phi <i>FEE, or FI (as in hi)</i>
κ	kappa <i>KAP-uh</i>	χ	chi <i>KI (as in hi)</i>
λ, Λ	lambda <i>LAM-duh</i>	ψ, Ψ	psi <i>SIGH, or PSIGH</i>
μ	mu <i>MEW</i>	ω, Ω	omega <i>oh-MAY-guh</i>

Capitals shown are the ones that differ from Roman capitals.

Alphabetical Index

preface, v